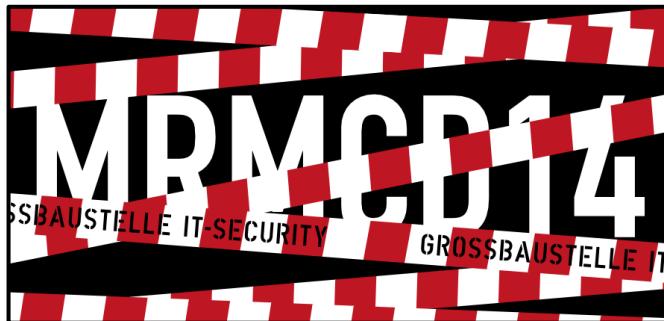


SCHIZO PHRENIC

FACADES



This talk is a collaboration with:

GYNAEEL COLDWIND

SECURITY RESEARCHER, GOOGLE



DRAGON SECTOR CAPTAIN

LIKES HAMBURGERS

[HTTP://GYNAEEL.COLDWIND.PL/](http://GYNAEEL.COLDWIND.PL/)



ANGE ALBERTINI

reverse engineering

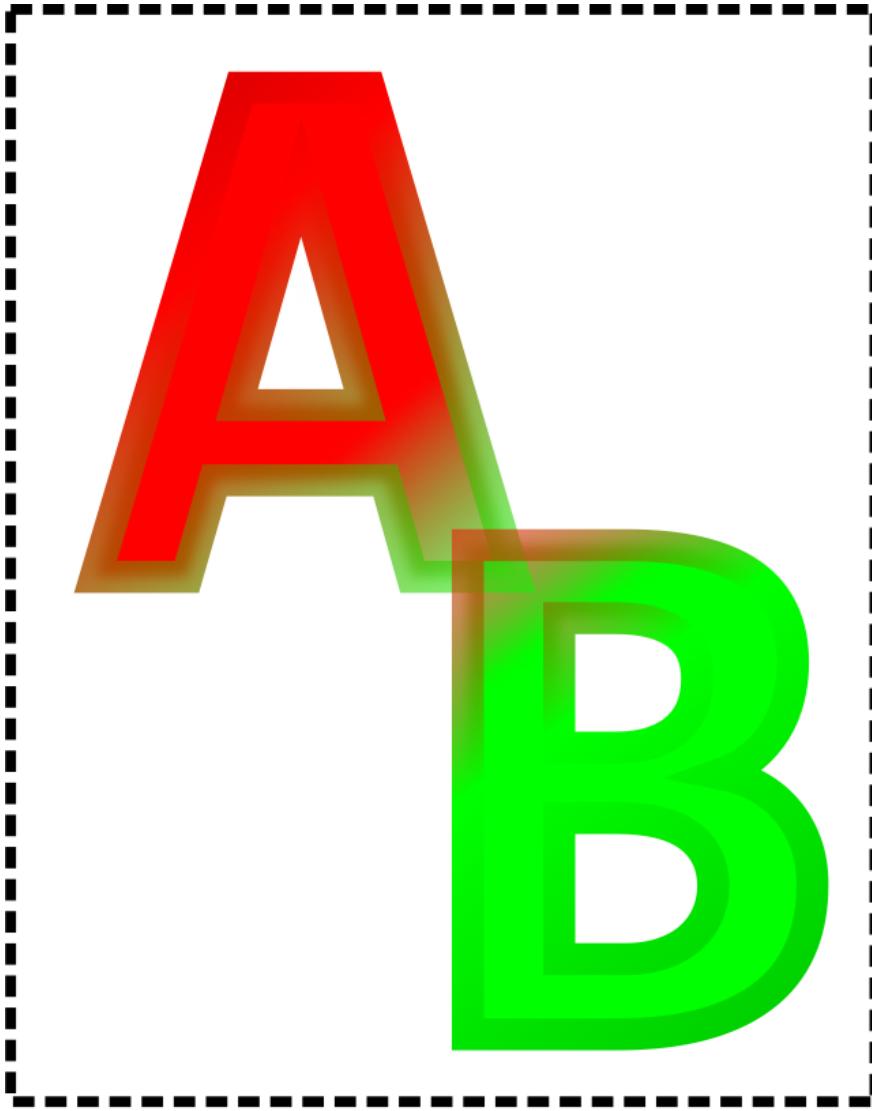
VISUAL DOCUMENTATIONS

@angealbertini

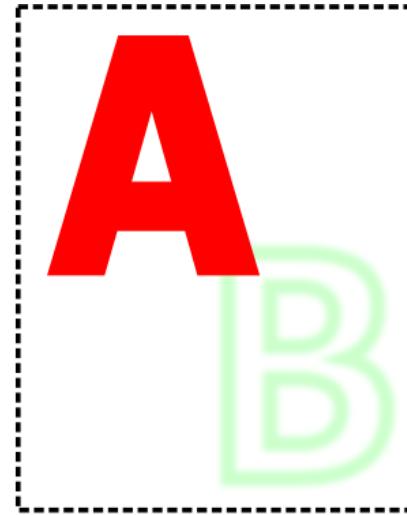
ange@corkami.com

<http://www.corkami.com>

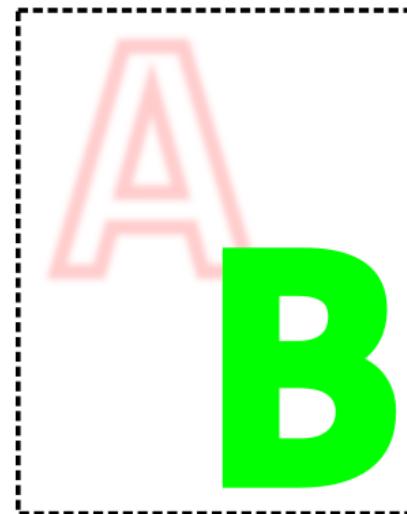




PQLG



Prg 1



Prg 2

**1 file, 2 programs
⇒ 2 different contents**

No active detection of the program in the file

Fooling, not failing

Both programs will load the file correctly:
No reported warning or error, no exploitation.

Abusing parsers for

- fun
- bypassing security
 - same-origin policy
 - evade detection
 - exfiltration
 - signing
 - Android Master Key

ZIP

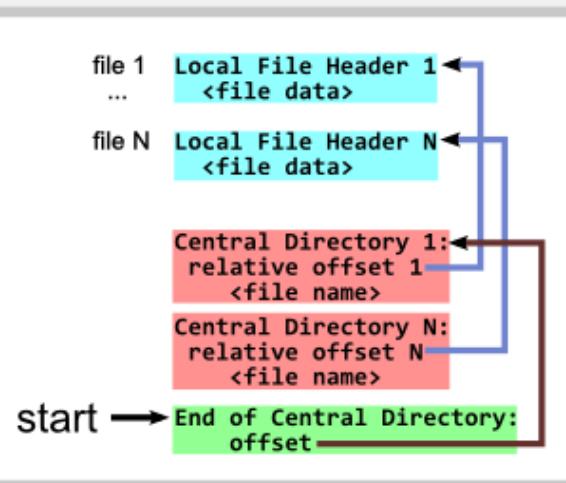


excerpt from Gynvael's talk:
"Dziesięć tysięcy pułapek: ZIP, RAR, etc."
(<http://gynvael.coldwind.pl/?id=523>)

```
~$ unzip simple.zip
Archive: simple.zip
extracting: hello.txt
~$ cat hello.txt
Hello World!
```

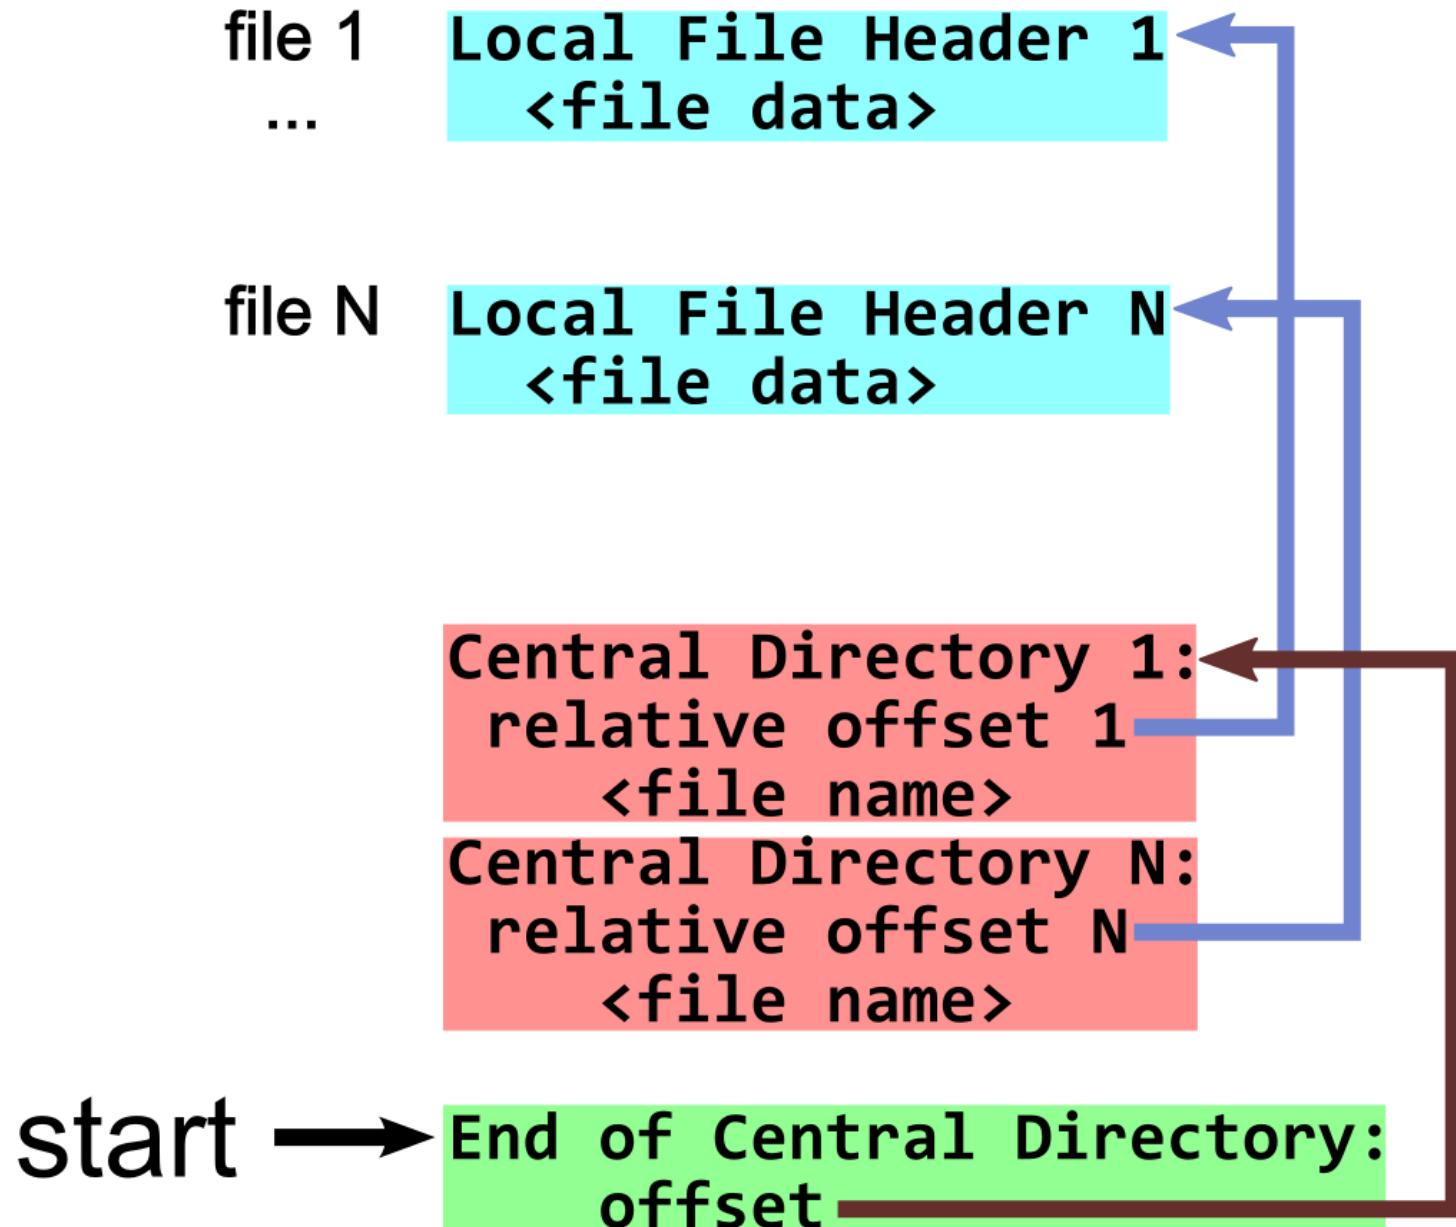
	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F		
00:	.P	.K	03	04	0A	00			00	00			DD	DD				
10:	14	7D	0D	00	00	00	0D	00	00	00	.	H	e					
20:	.l	.l	o	.	W	.o	r	l	d	!	0A	.P	.K	01	02			
30:	0A	00									DD	DD	14	7D	0D			
40:	00	00	00	0D	00	00	00	09	00									
50:										00	00	00	00	h	e			
60:	x	t	.	P	.K	05	06			01	00	37	00	l	l	o	..	t
70:	00	00	2B	00	00	00	00											

SIMPLE.ZIP

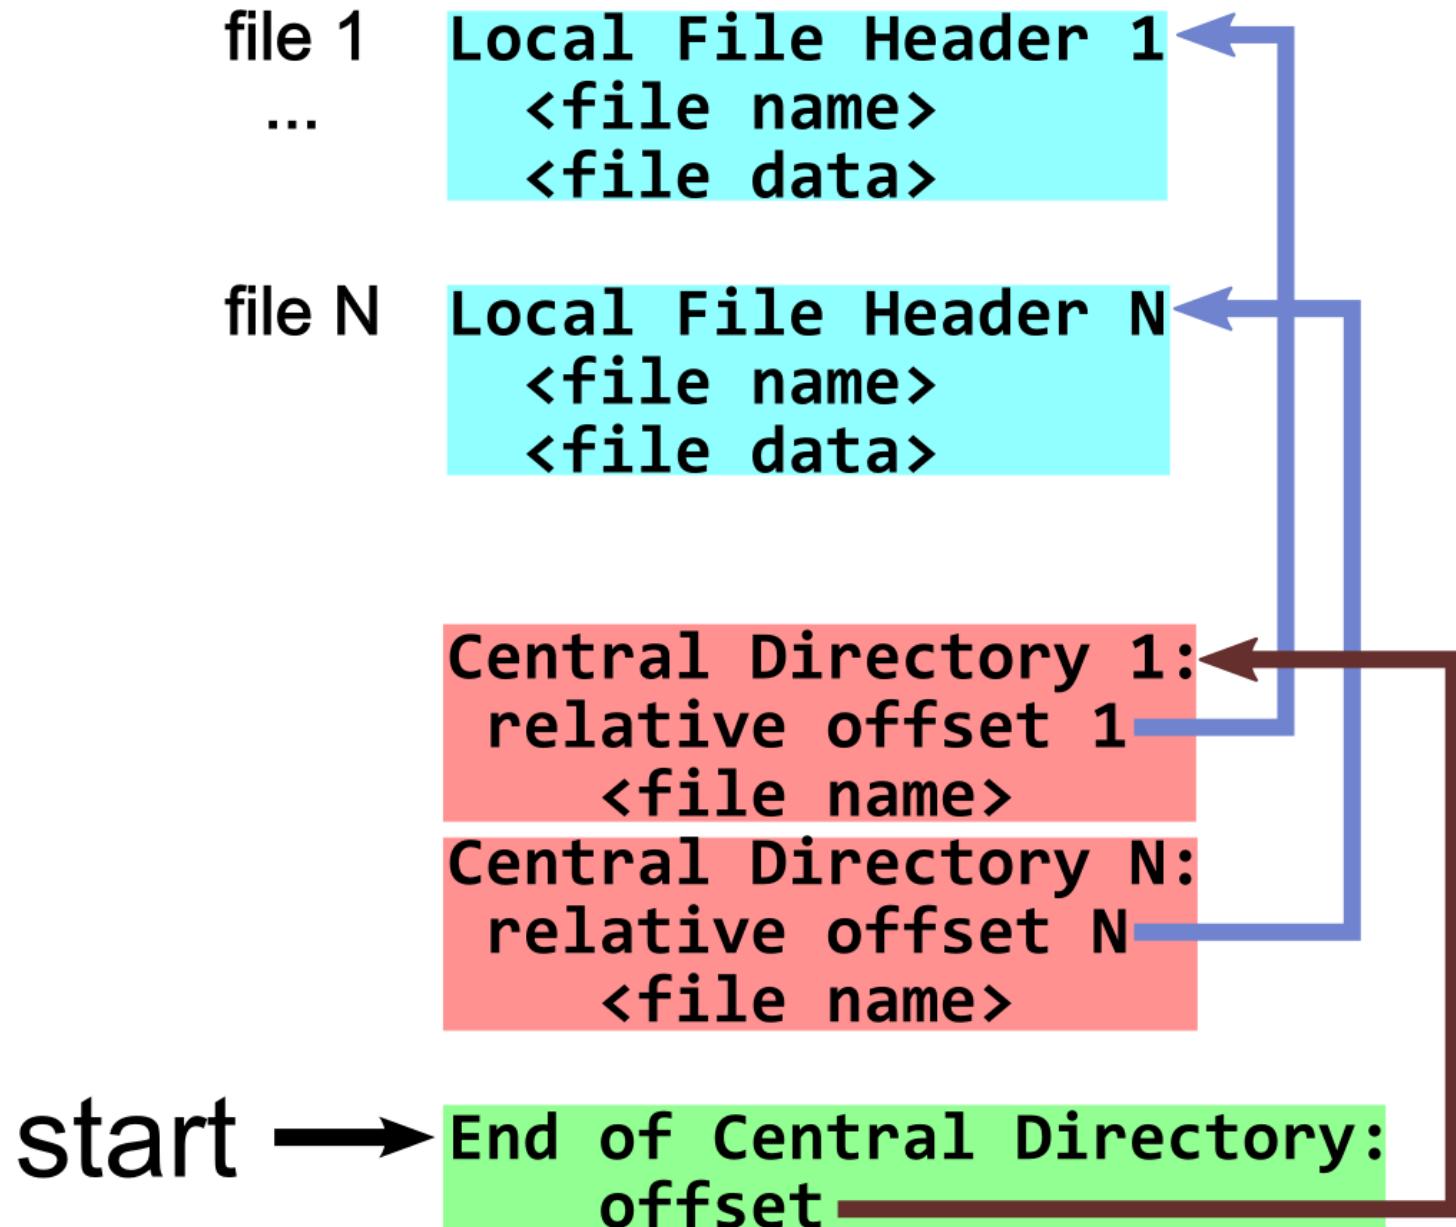


	description	value
Local File Header	local file header signature version needed to extract compression method crc-32 compressed size uncompressed size	PK\x03\x04 10 (default value) 0 (no compression) 0x7D14DDDD 0x0D 0x0D
file data	file data	Hello World!\n
Central Directory	central file header signature version needed to extract crc-32 compressed size uncompressed size file name length relative offset of local header 0	PK\x01\x02 10 (default value) 0x7D14DDDD 0x0D 0x0D 9 0 →
file name	file name	hello.txt
End of Central Directory	end of central dir signature total number of entries in the central directory size of the central directory offset of start of central directory with respect to the starting disk number	PK\x05\x06 1 0x37 0x2B →

ZIP ARCHIVE

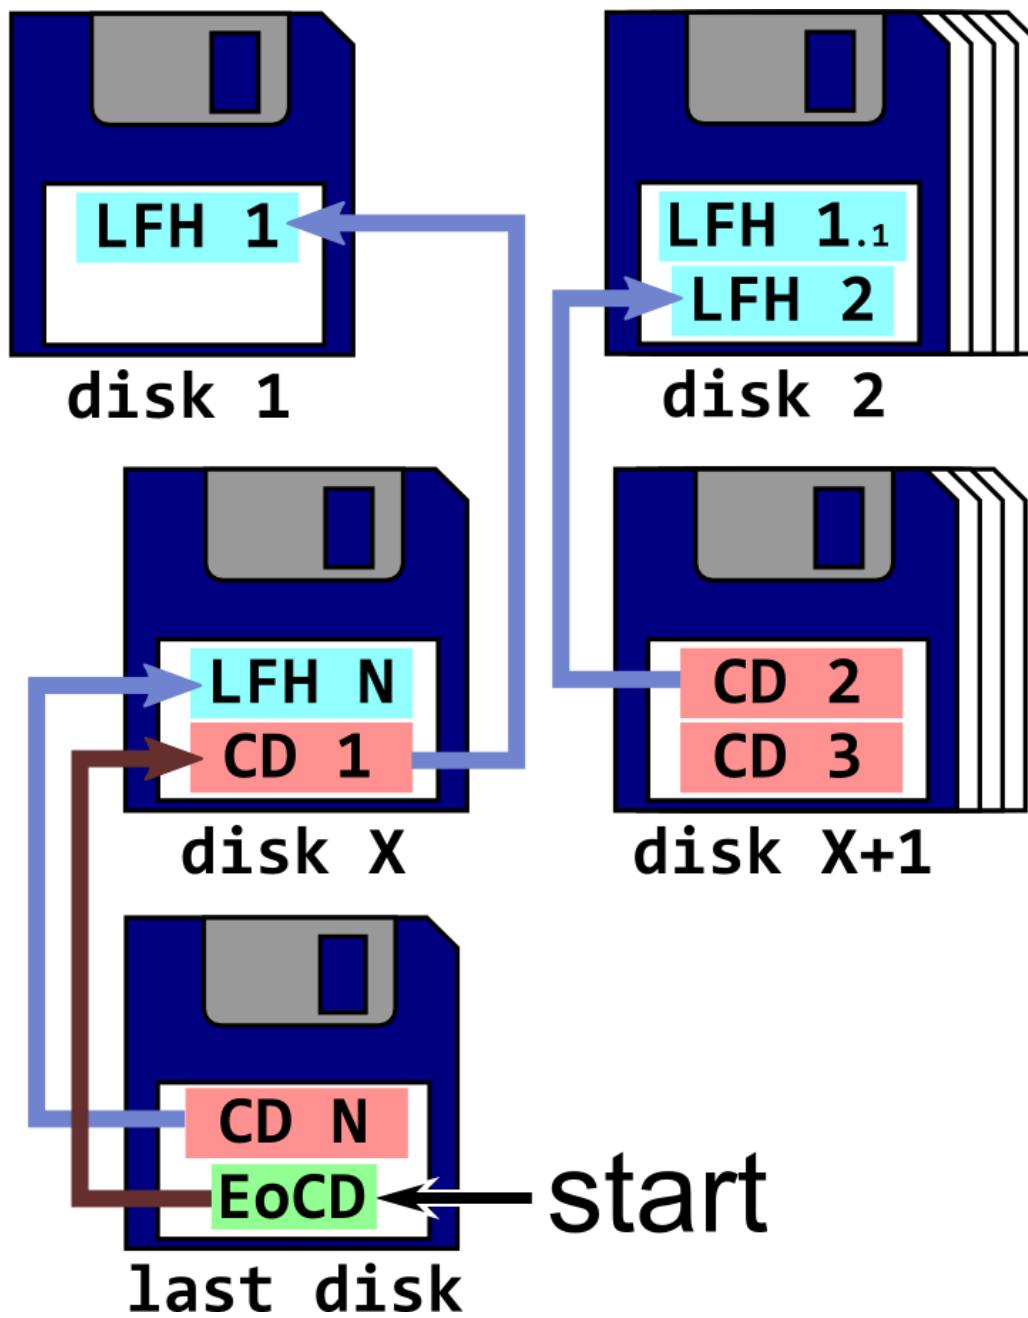


ZIP structures are parsed from the end.



File names are actually duplicated.

Why this weird structure?



ZIP archives were commonly read & written on the fly over multiple floppies.

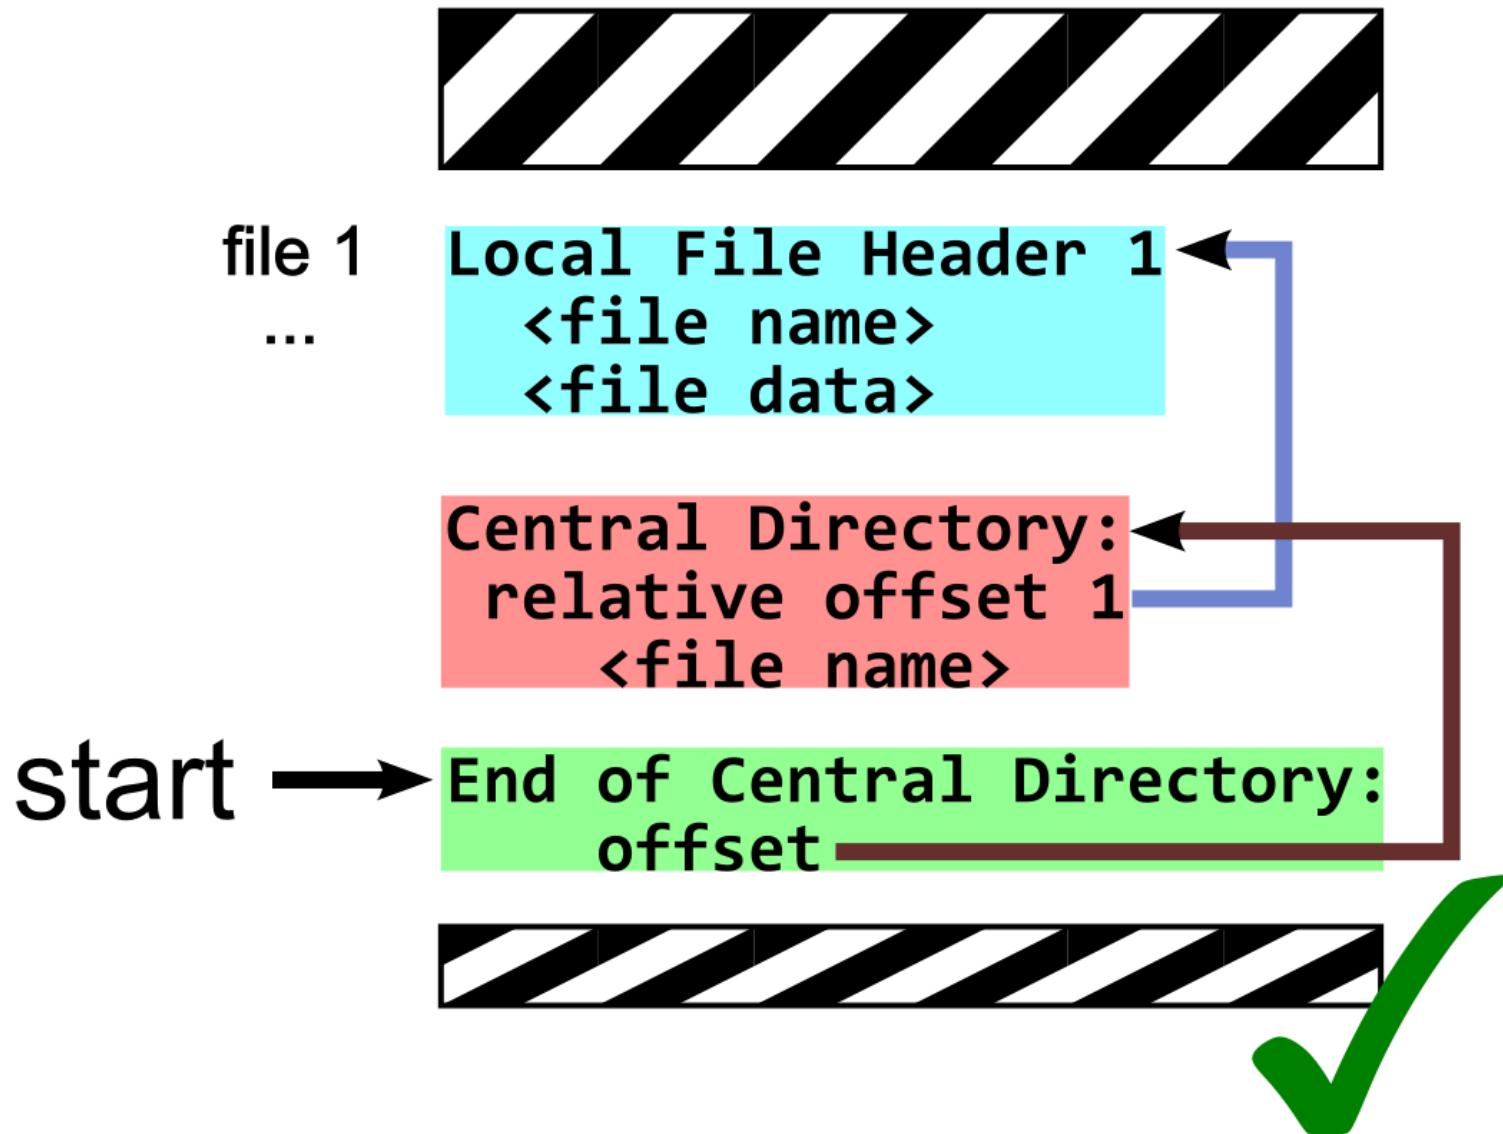
Minimize floppy swaps

- creation
 - a. create one LFH per file
 - Floppy full ⇒ start a new LFH on the next floppy
 - b. when all files are finished, write CDs sequence (1/file)
 - c. when all CDs are written, write the EoCD
- extraction
 - a. insert last floppy (contains the EoCD)
 - b. insert the floppy with 1st CD
 - (often, the last floppy contains EoCD + all CDs)
 - c. insert the corresponding LFH's first floppy
 - insert next floppies if required

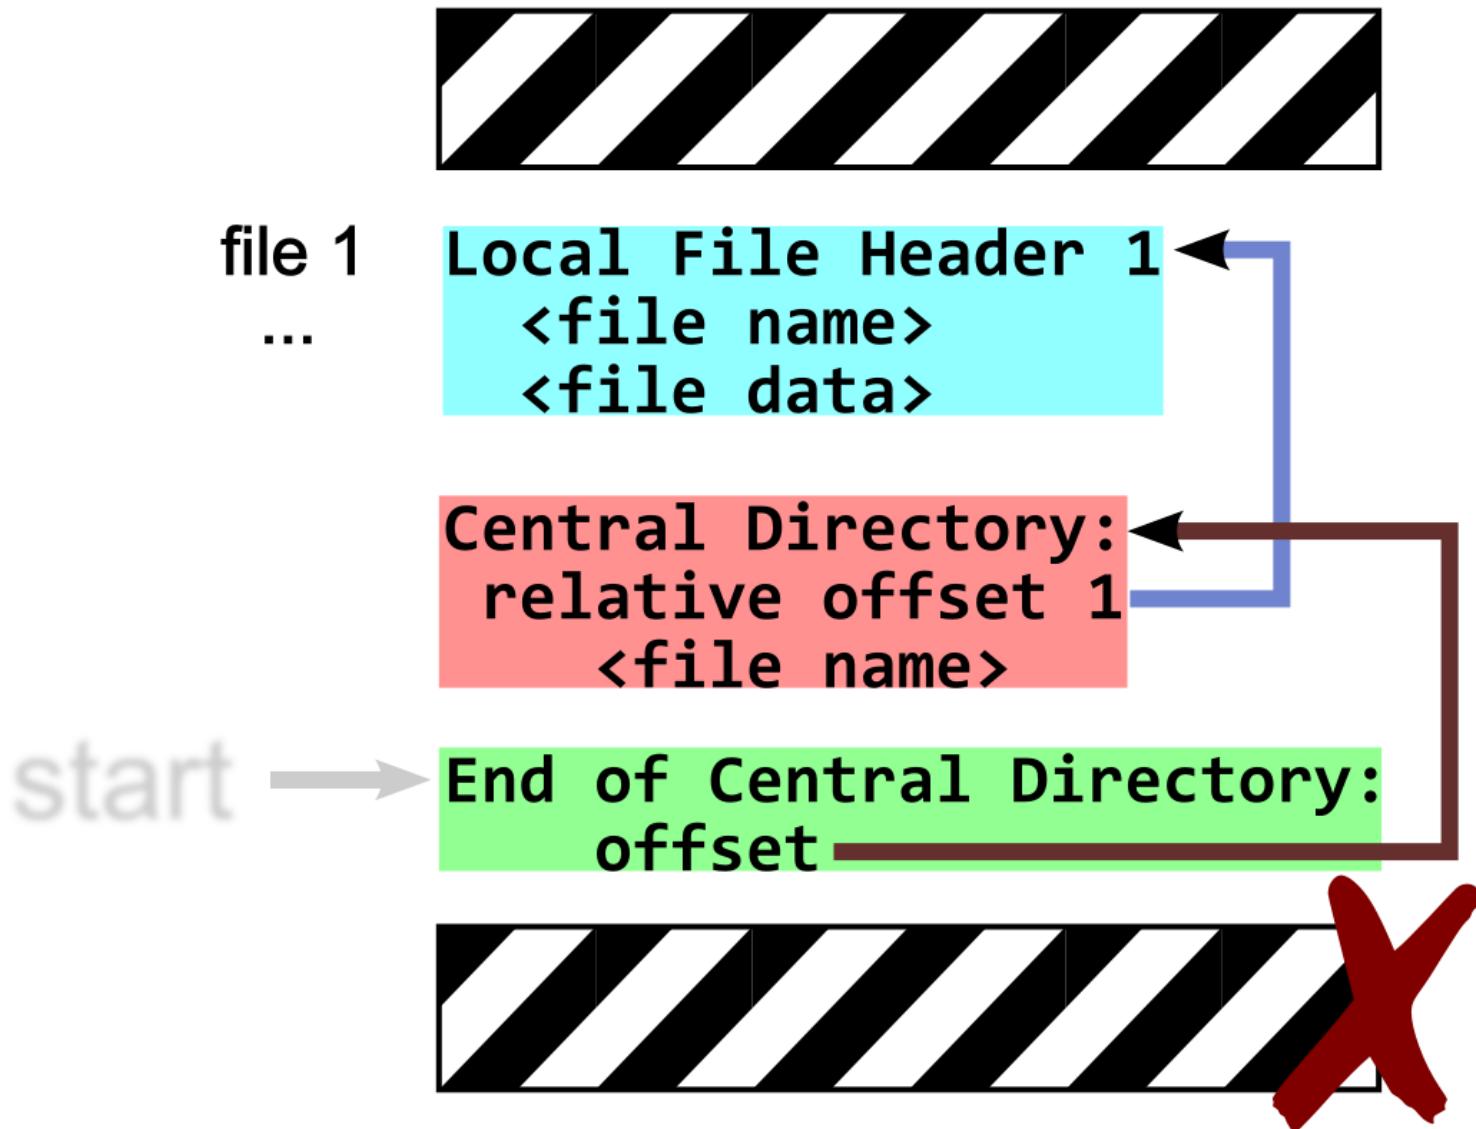
**ZIP was very useful,
but *now* it's awkward.**

Newer archive formats are parsed top-down.

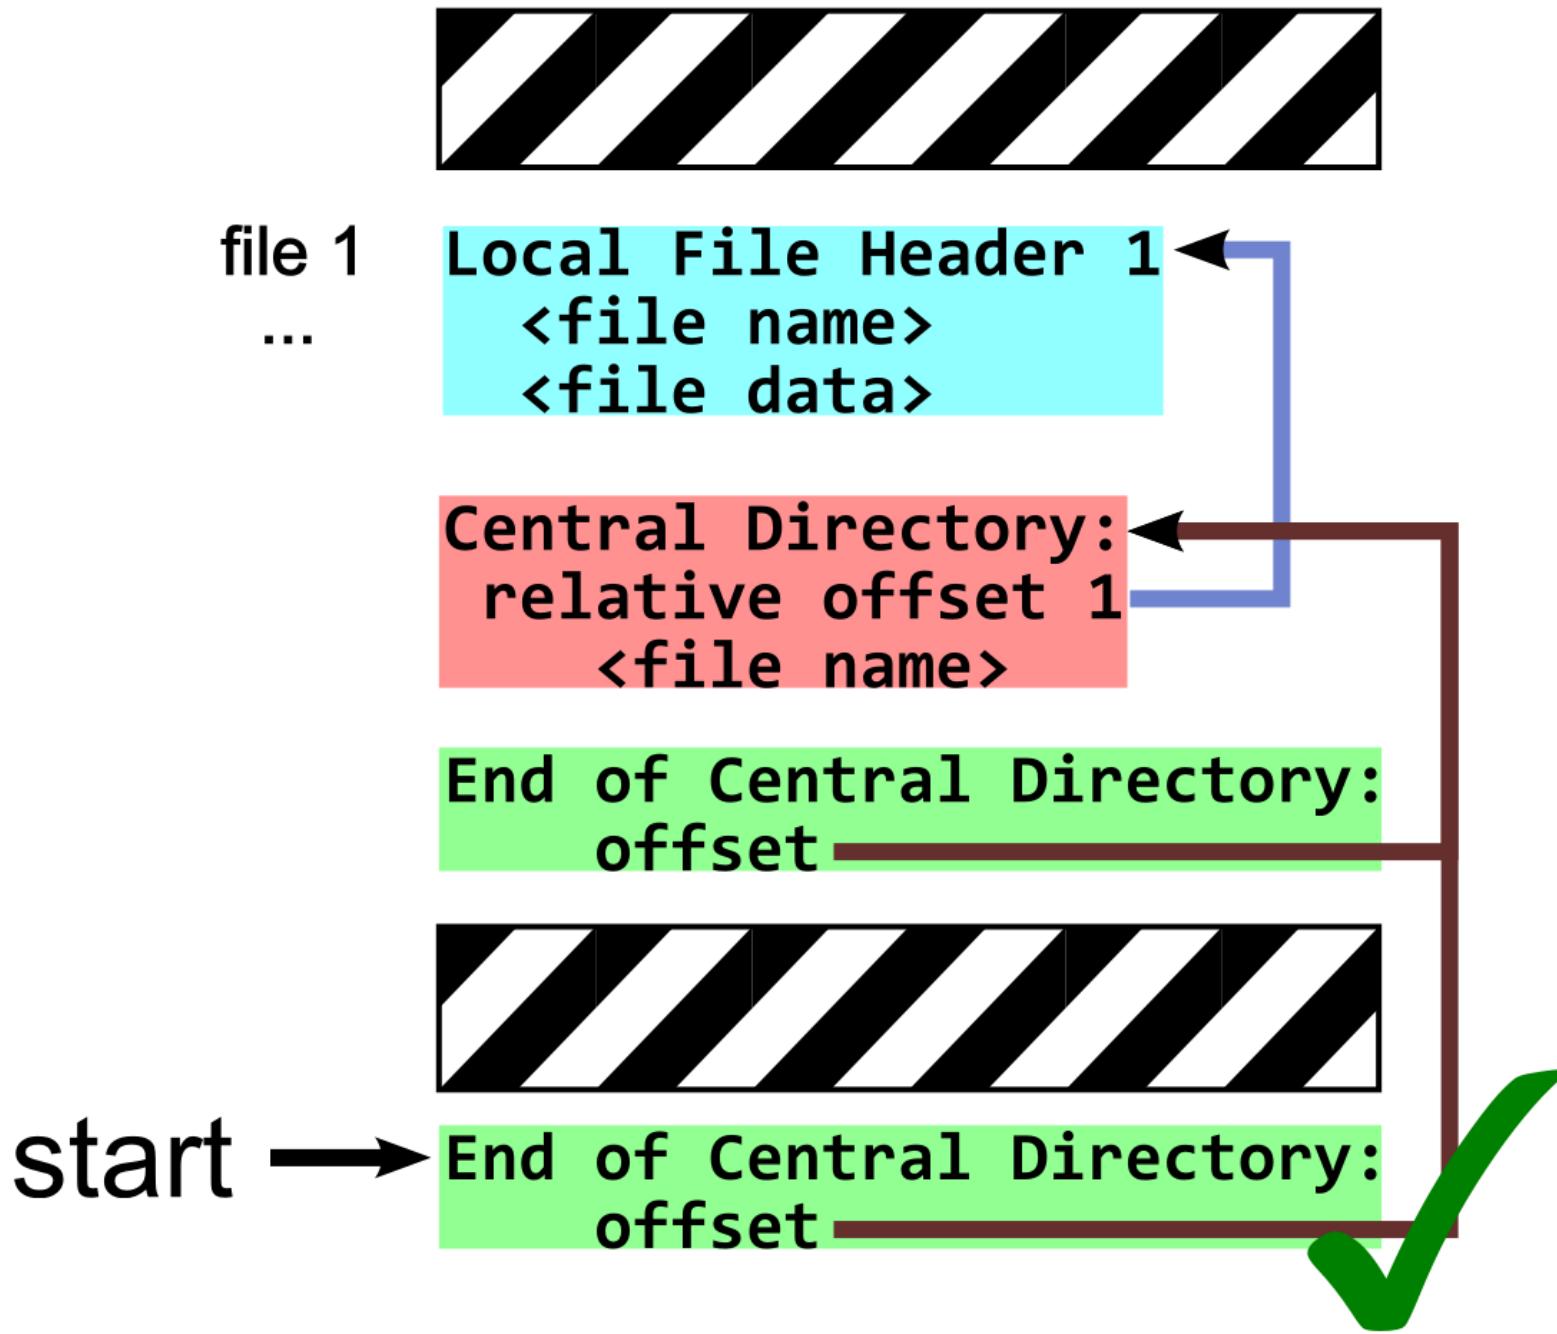
Position in the file



Prepended and appended data is tolerated...

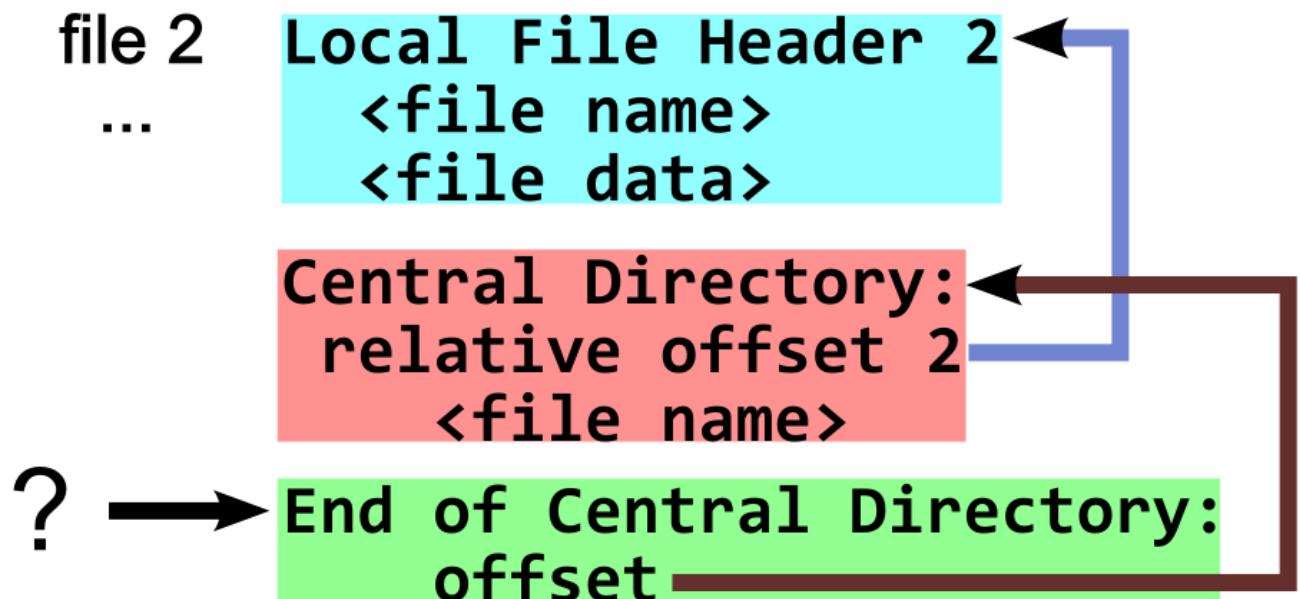
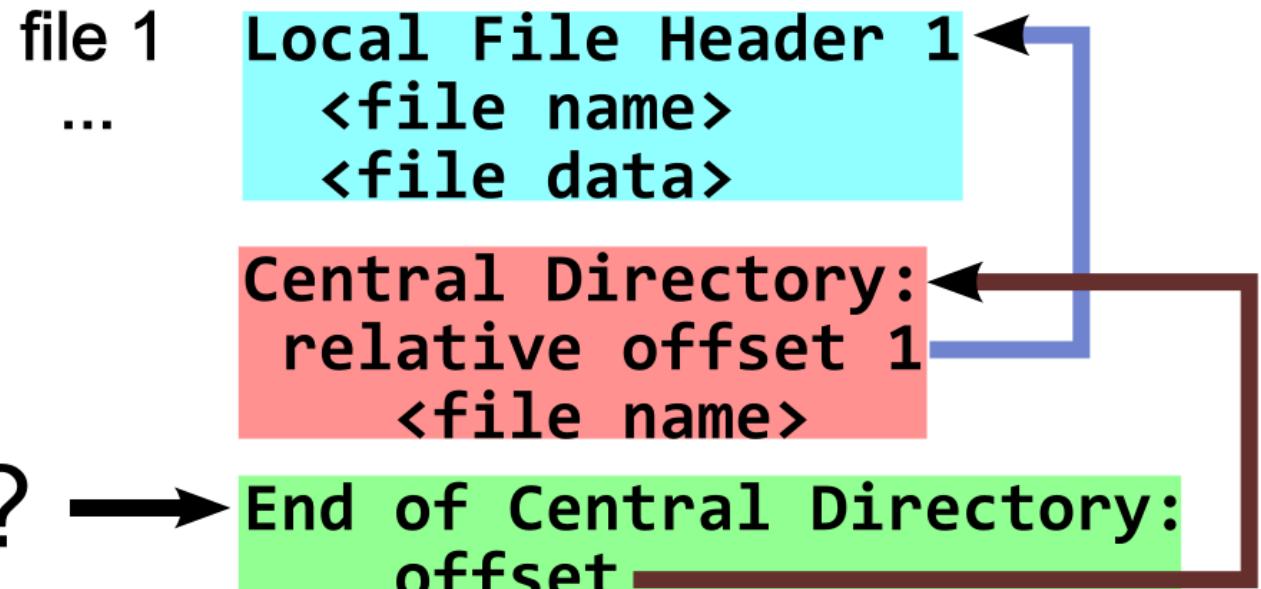


...but not too much!
(for obvious performance reason)

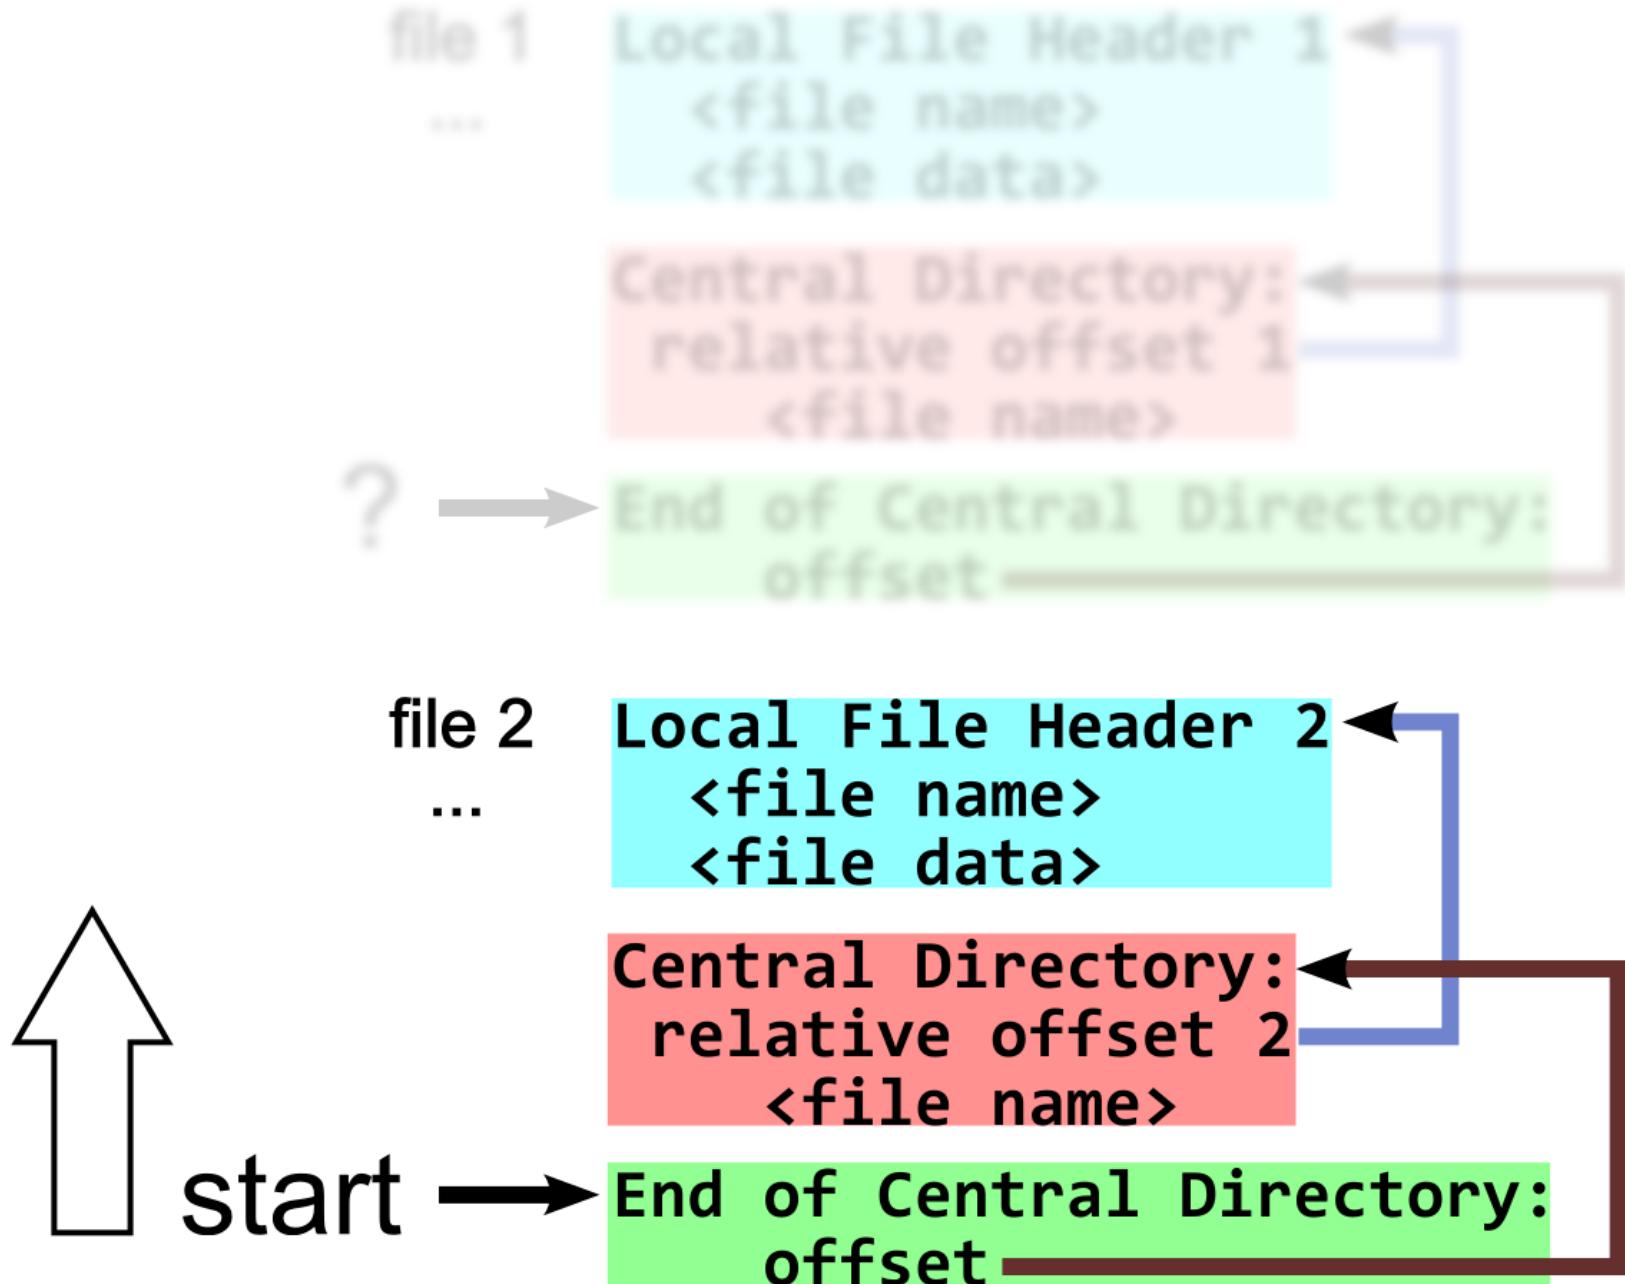


Duplicating the (relatively small) EoCD increases compatibility.

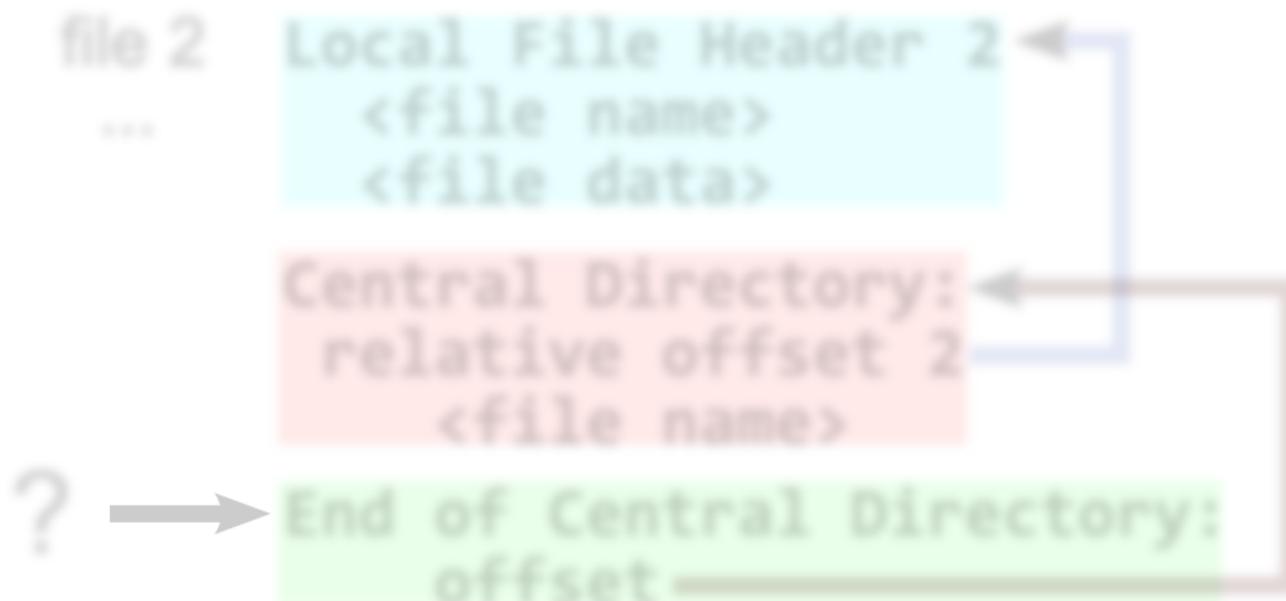
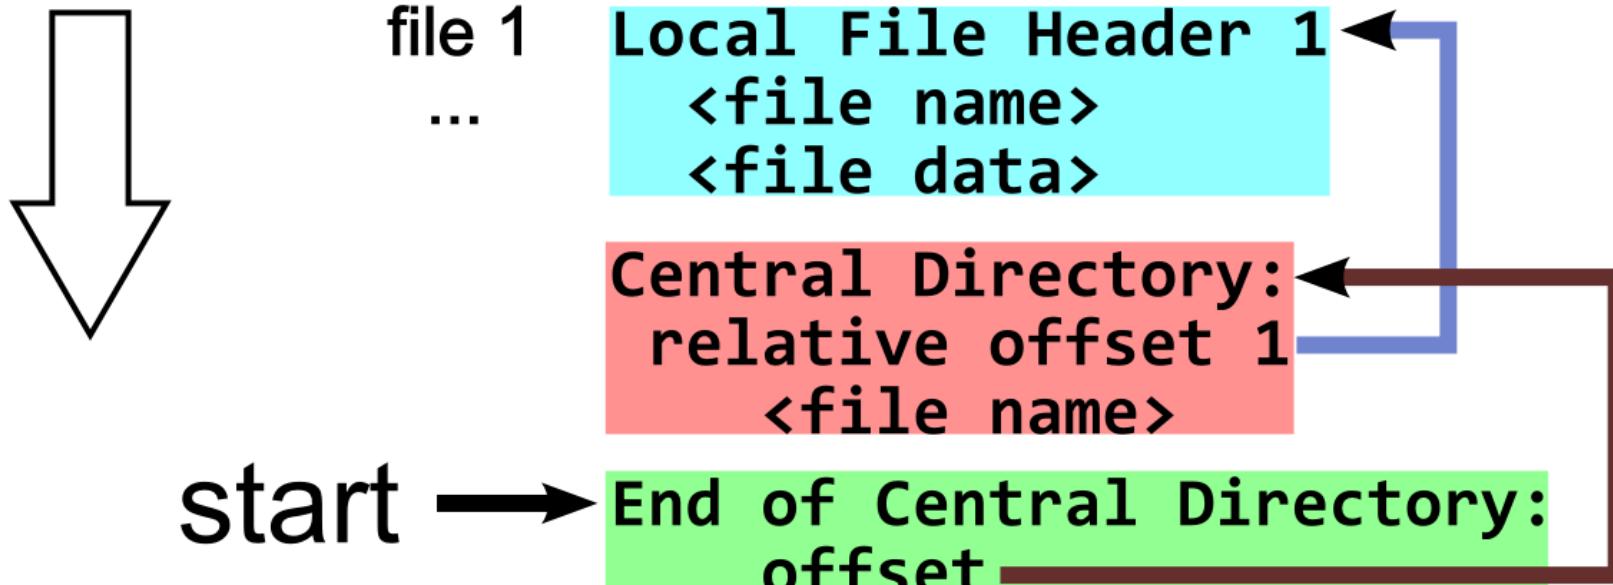
Scanning direction



If you concatenate 2 archives...

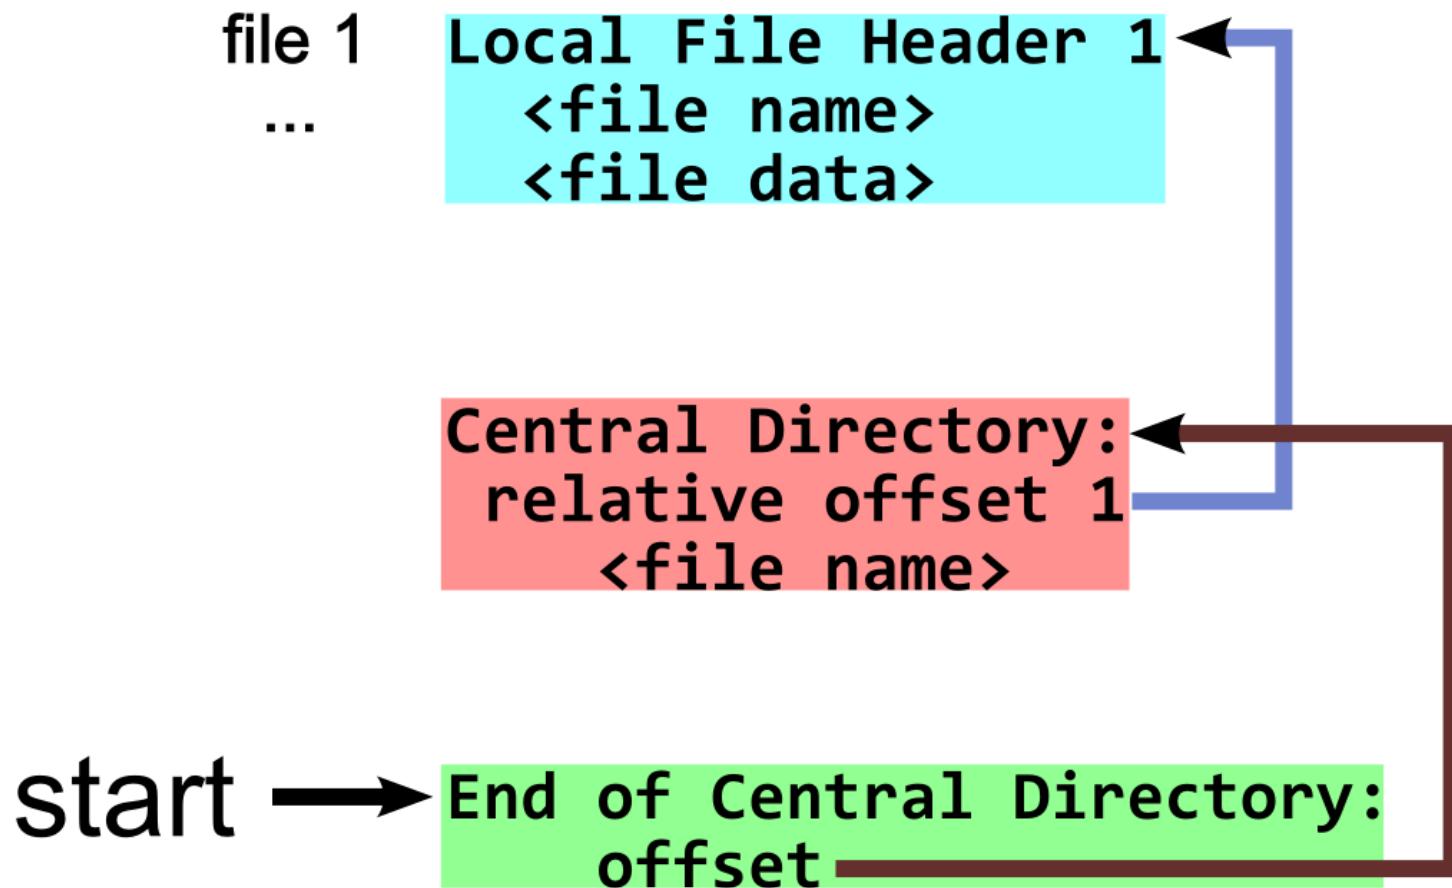


...if you parse bottom-up (standard), you find the 2nd one...



...but you will get the other archive if you parse top-down.

Superfluous headers



You could parse everything nicely...

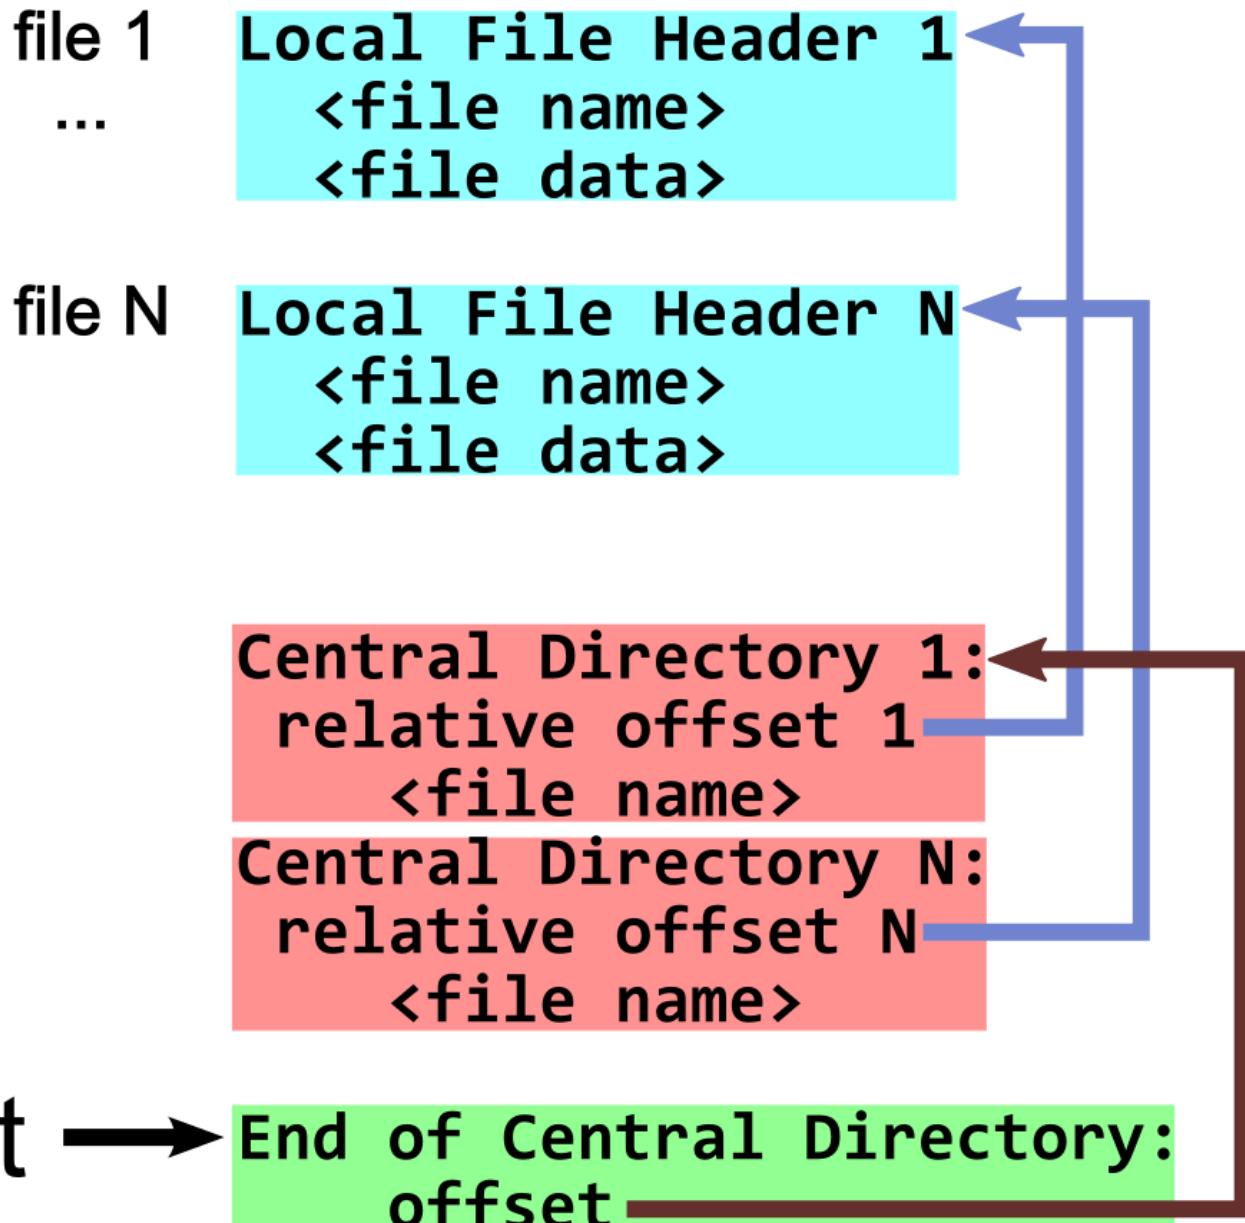
file 1 Local File Header 1
... <file name>
 <file data>

Central Directory:
relative offset 1
 <file name>

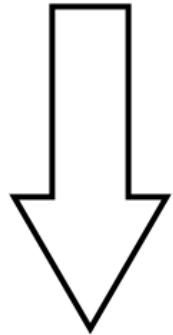
start → End of Central Directory
 offset nnnn

...but in the end, only the Local File Headers matter.

1 file = 1 Local File Header



Since most ZIP archives start with a sequence of Local File Headers...



file 1

...

Local File Header 1
 <file name>
 <file data>

file N

Local File Header N
 <file name>
 <file data>

Central Directory:
relative offset 1
 <file name>

Central Directory:
relative offset 8
 <file name>

start →

End of Central Directory:
offset n

You can parse them top-down (until a break) and ignore the CD and EoCD.

**Standard parsing:
bottom-up + all headers**

“Efficient” parsing: top-down + LFHs only

Not standard, but good enough in most cases.

**Nowadays, most ZIPs are
a sequence of LFHs
from the start**

ZIP Archive comment



**End of Central Directory:
offset**

comment length

comment:

The EoCD contains an optional comment field...



End of Central Directory:

offset

comment length

comment:

Local File Header 2

<file name>

<file data>

Central Directory:

relative offset 2

<file name>

End of Central Directory:

offset

...that can contain a complete archive !

Recap

- Parsing direction:
 - standard is bottom-up
 - parsing LFHs from the start would work in most cases
- ZIP should be located near the end of the file
 - or at least, its EoCD
- An archive comment can contain another complete archive

Let's test the parsers!

abstract.zip



LFH S

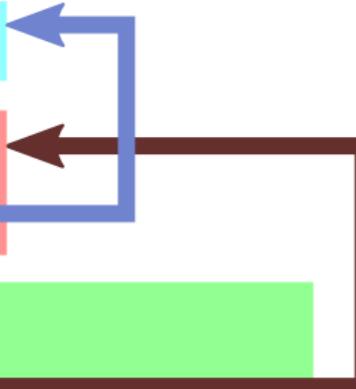


LFH A

LFH D

**CD D:
rel.offset**

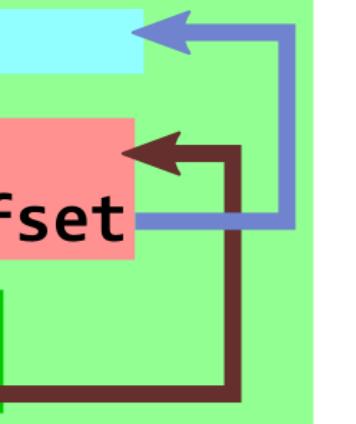
**EoCD D:
offset**



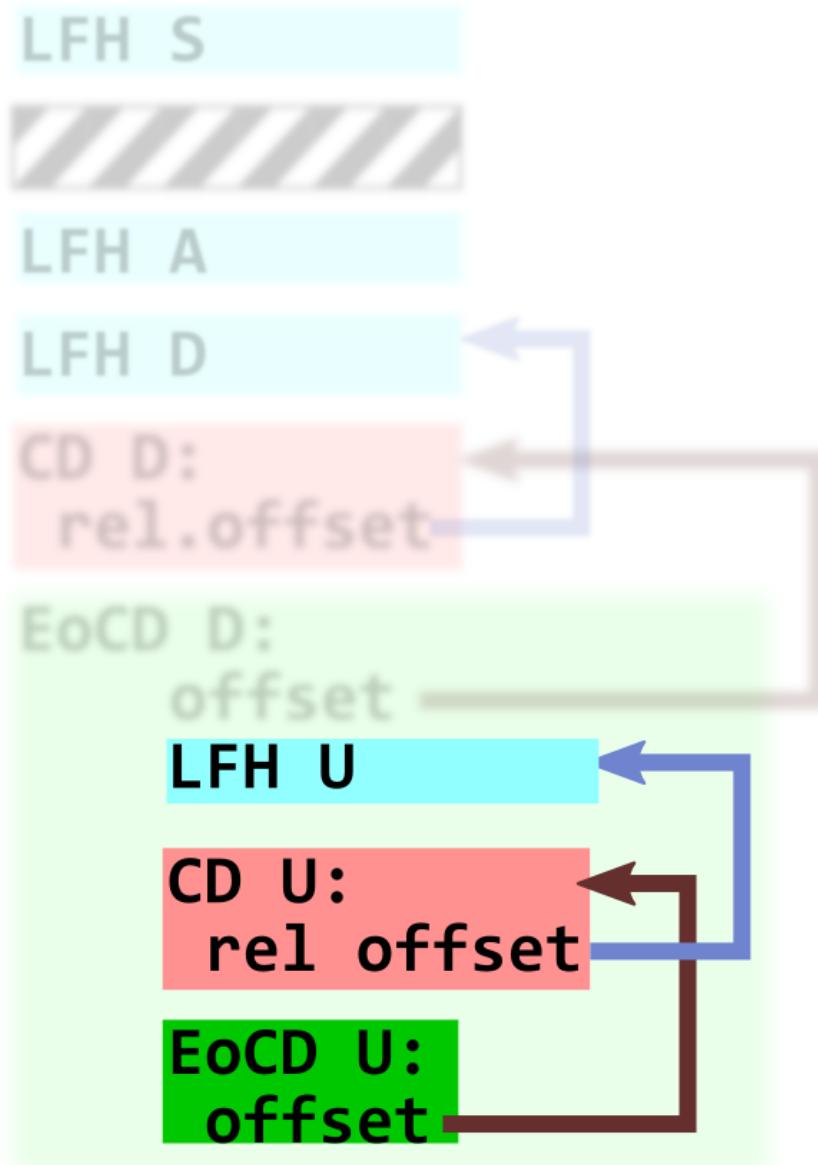
LFH U

**CD U:
rel offset**

**EoCD U:
offset**



4 LFHs, 4 ways to parse this archive:



7-zip
C# DotNetZip
Java ZipFile
OSX UnZip
perl Archive::Zip
Python zipfile
Total Commander 8.01
UnZip 6.00 (Debian)
Windows 7 Explorer
WinRAR
Ada Zip-Ada v45
ALZip
b1.org
Midnight Commander
KGB Archiver
JSZip
Jeffrey's Exif Viewer
WOZIP
GNOME File Roller
zip.vim v25
Emacs Zip-Archive mode
Go archive/zip
Pharo smalltalk 2.0 ZipArchive
Ubuntu less

1/ you parse it bottom-up



PHP ZipArchive
PHP zip_open ...
PHP zip:// wrapper
tcl + tclvfs + tclunzip

2/ you parse it top-down

LFH S



LFH A

LFH D

CD D:
rel.offset

EoCD D:

offset

LFH U

CD U:

rel offset

EoCD U:

offset

Ruby **rubyzip2**

Java **ZipArchiveInputStream**

java.util.zip.ZipInputStream

3/ look for LFHs from the start (until a break)

LFH S

binwalk

LFH A

LFH D

CD D:

rel.offset

EoCD D:

offset

LFH U

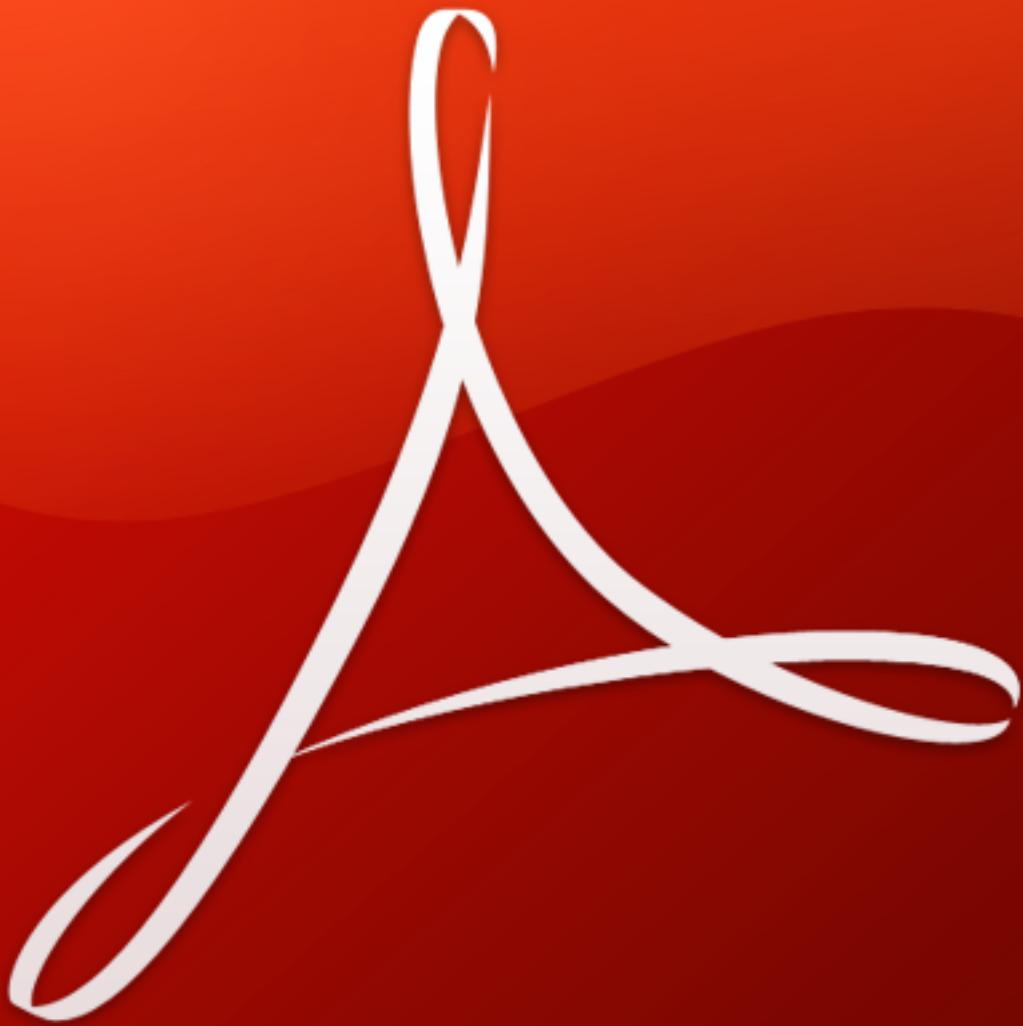
CD U:

rel offset

EoCD U:

offset

4/ scan for LFHs aggressively (you get all four)



PDF 101

basics of the PDF file format

Part II / III



<http://youtu.be/JQrBgVRgqtc?t=11m15s>

<https://speakerdeck.com/ange/pdf-secrets-hiding-and-revealing-secrets-in-pdf-documents?slide=44>

```
%PDF-1.1
```

```
1 0 obj
<<
/Pages 2 0 R
>>
endobj

2 0 obj
<<
/Type /Pages
/Count 1
/Kids [3 0 R]
>>
endobj

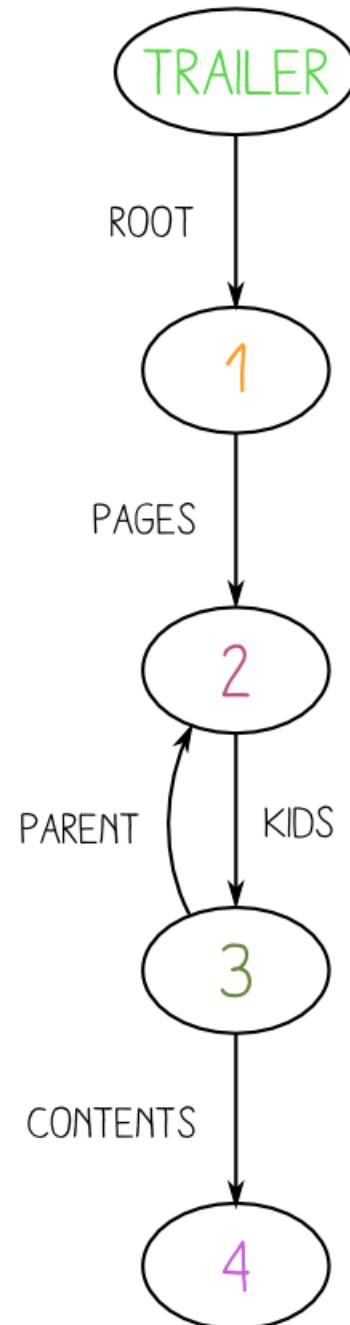
3 0 obj
<<
/Type /Page
/Contents 4 0 R
/Parent 2 0 R
/Resources <<
/Font <<
/F1 <<
/Type /Font
/Subtype /Type1
/BaseFont /Arial
>>
>>
>>
endobj

4 0 obj
<< /Length 50 >>
stream
BT
/F1 110 Tf
10 400 Td
(Hello World!)Tj
ET
endstream
endobj

xref
0 5
0000000000 65535 f
0000000010 00000 n
0000000047 00000 n
0000000111 00000 n
0000000313 00000 n

trailer
<<
/Root 1 0 R
>>

startxref
413
%%EOF
```



PDF Trick #1

trailers

7.5.5 File Trailer

The *trailer* of a PDF file enables a conforming reader to quickly find the cross-reference table and certain special objects. Conforming readers should read a PDF file from its end. The last line of the file shall contain only the end-of-file marker, **%%EOF**. The two preceding lines shall contain, one per line and in order, the keyword **startxref** and the byte offset in the decoded stream from the beginning of the file to the beginning of the **xref** keyword in the last cross-reference section. The **startxref** line shall be preceded by the *trailer dictionary*, consisting of the keyword **trailer** followed by a series of key-value pairs enclosed in double angle brackets (`<<...>>`) (using LESS-THAN SIGNs (3Ch) and GREATER-THAN SIGNs (3Eh)). Thus, the trailer has the following overall structure:

```
trailer
  << key1 value1
      key2 value2
      ...
      keyn valuen
  >>
startxref
Byte_offset_of_last_cross-reference_section
%%EOF
```

trailer ⇒ root object ⇒ complete document

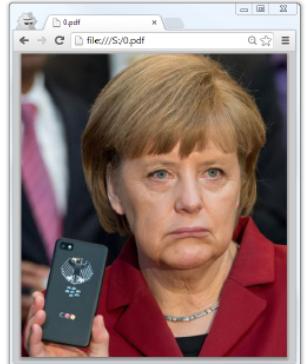
% trailer <</Root ...>>

trailer <</Root ...>>

<</Root ...>>

a line comment - a correct trailer - a corrupted trailer

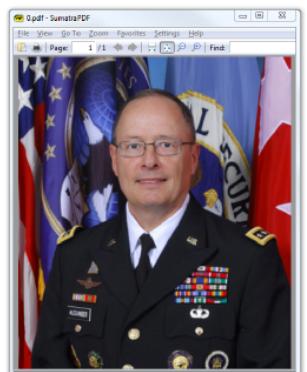
% trailer <</Root ...>>



trailer <</Root ...>>



<</Root ...>>



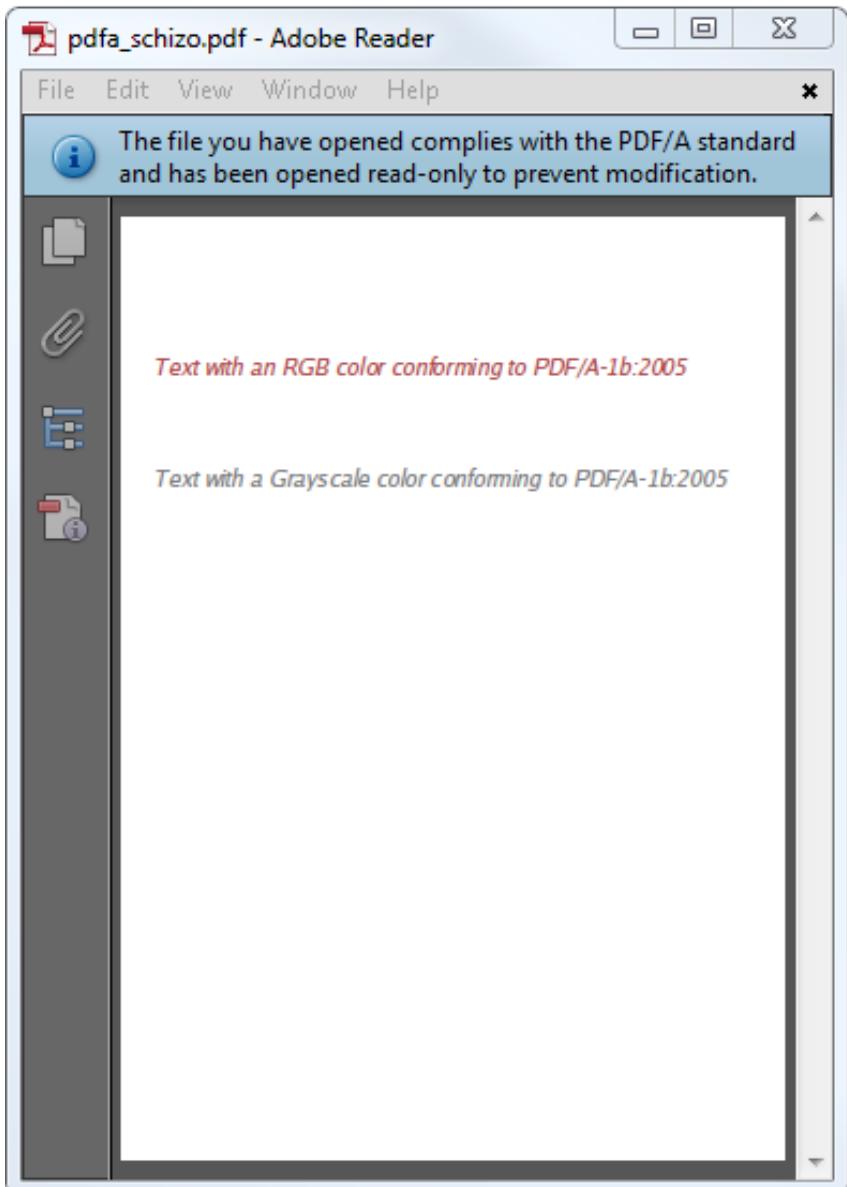
Each reader sees a different trailer.

PDF parsing

Each reader sees a completely different document

3 co-existing documents, all parsed through

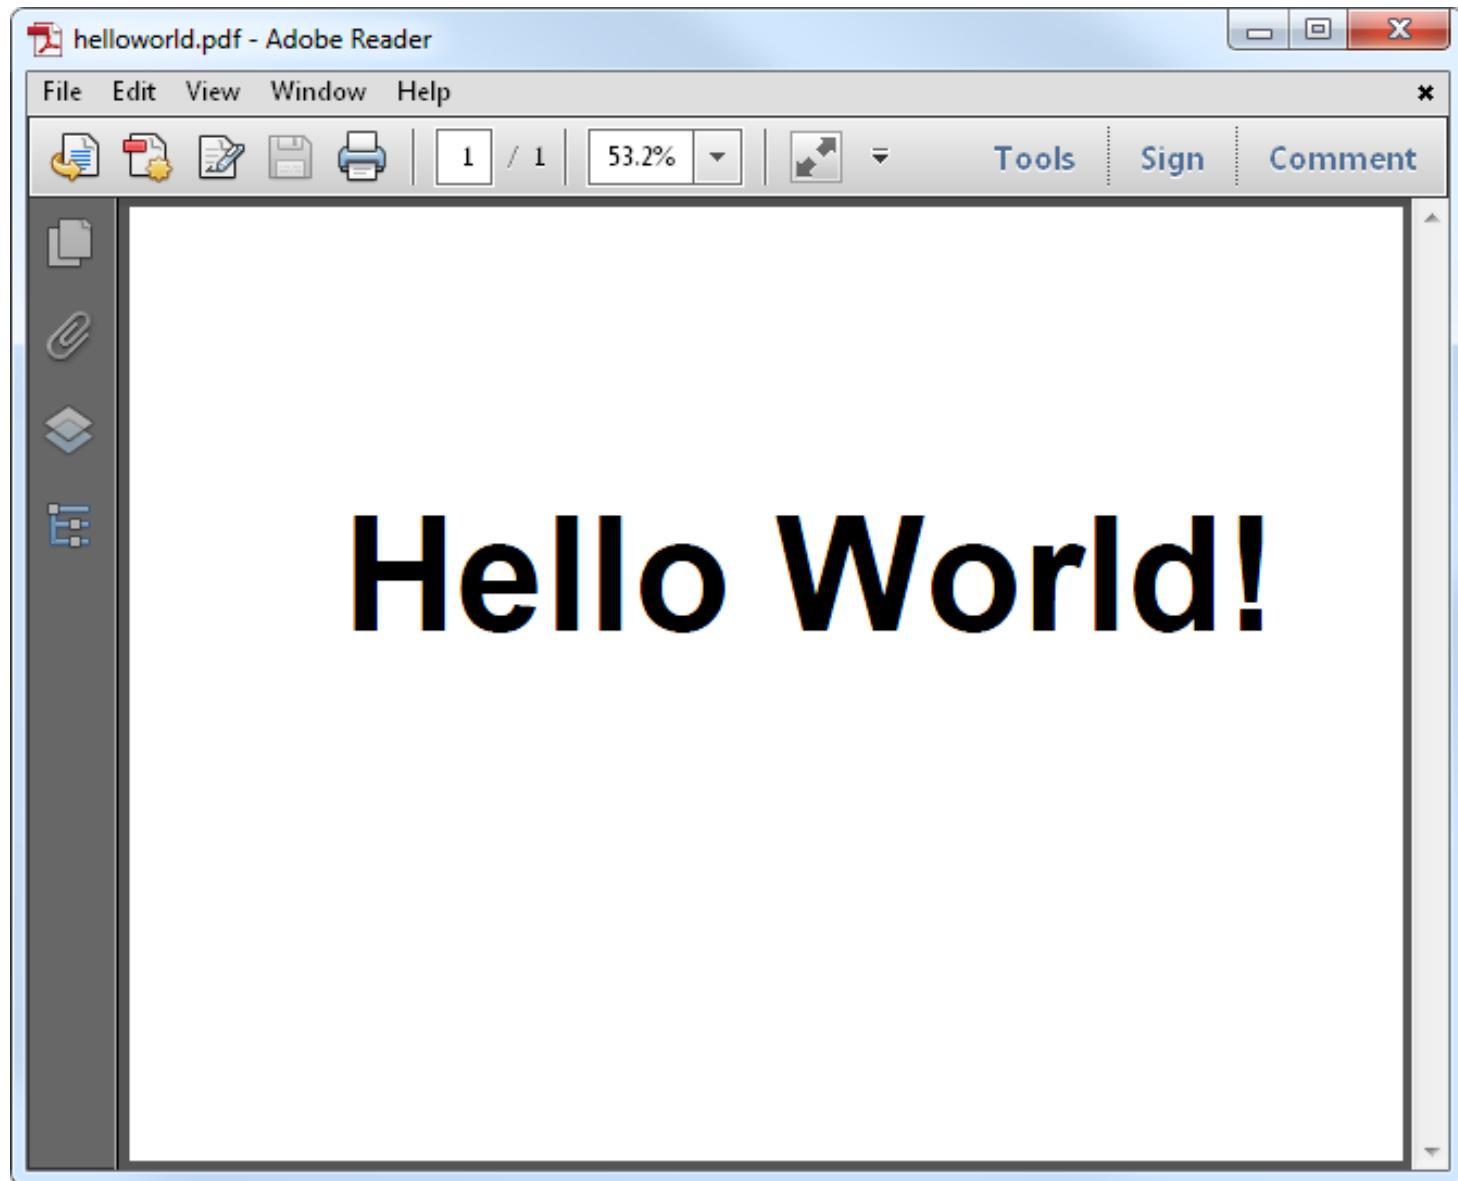
Viewers tolerance makes foreign elements ignored



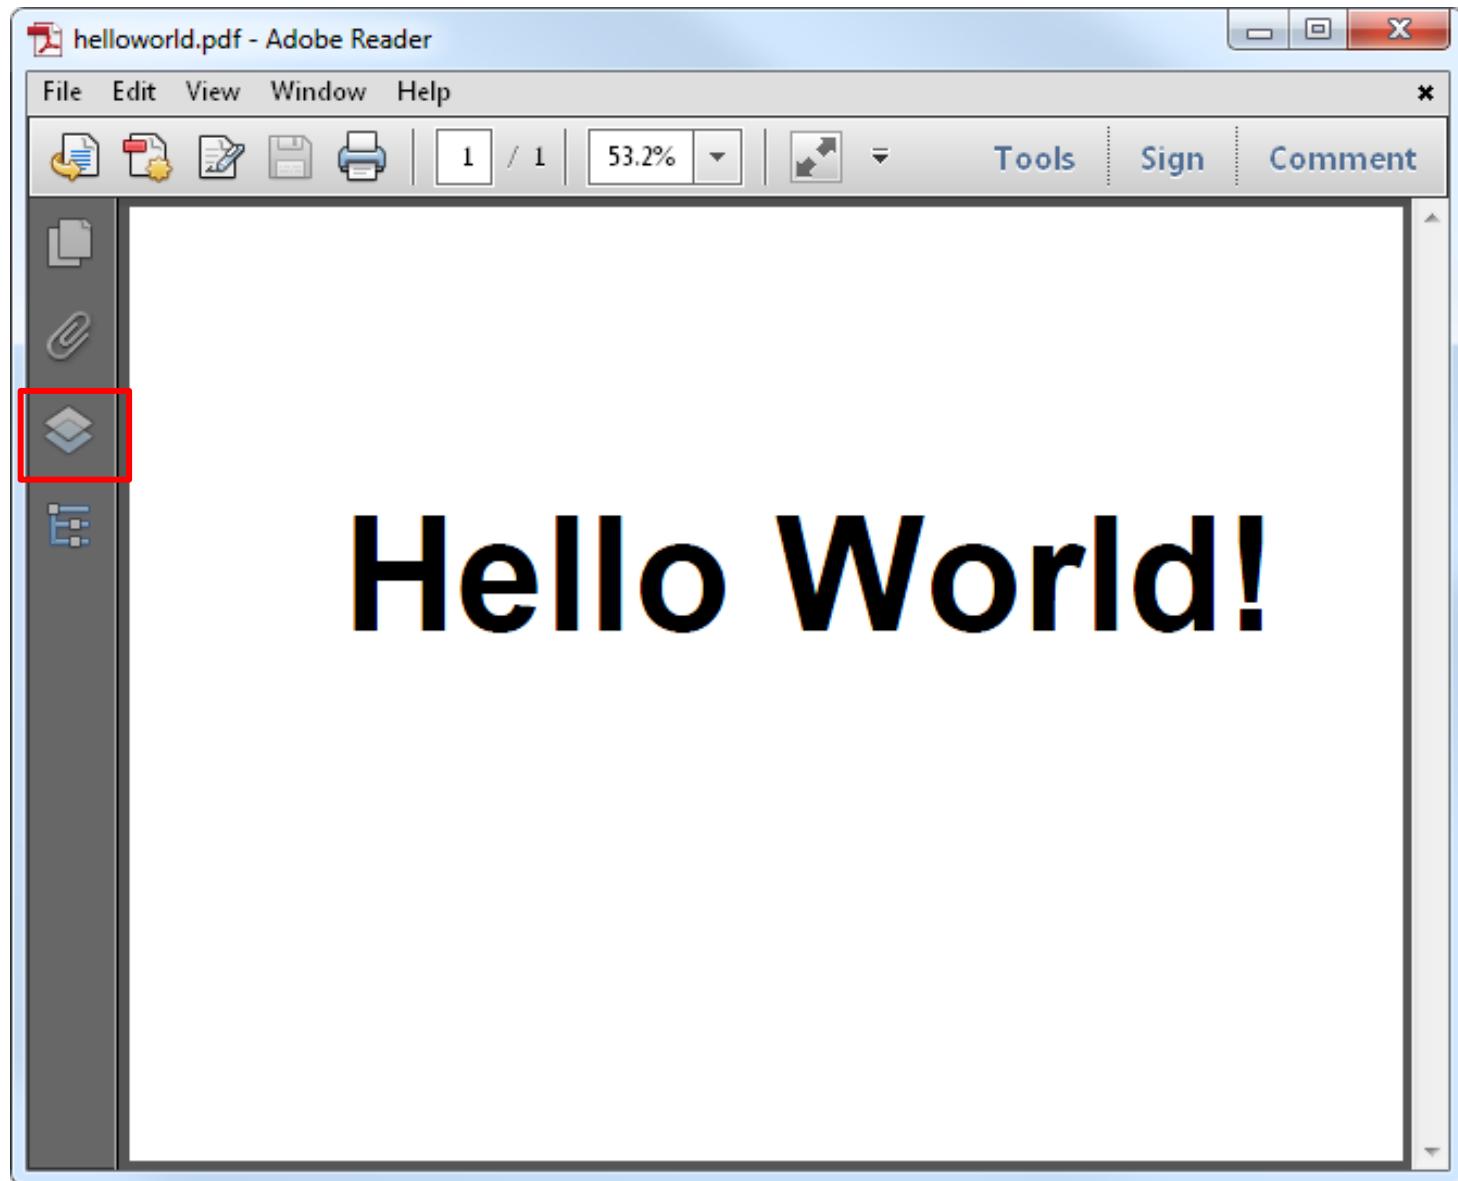
Also available in PDF/A flavor
(OK for Adobe Reader, but not for Preflight)

**sometimes,
it's in the specs...**

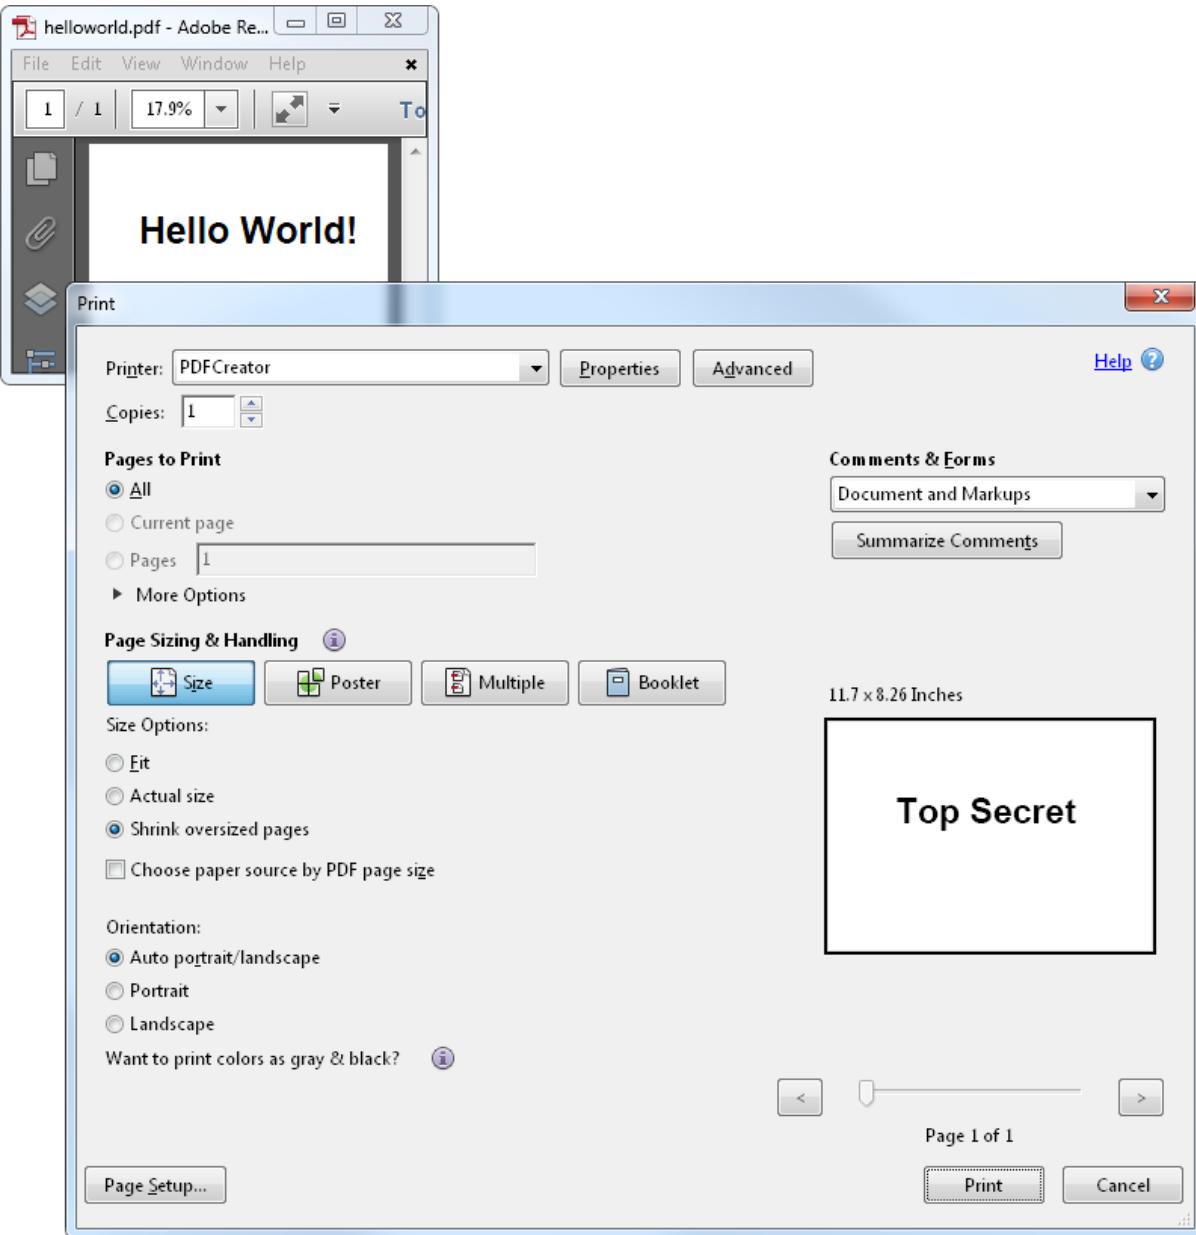
...but who knows all of them ?
(obscurity via over-specification)



Notice anything unusual?



This document contains layers.
(an advanced feature)

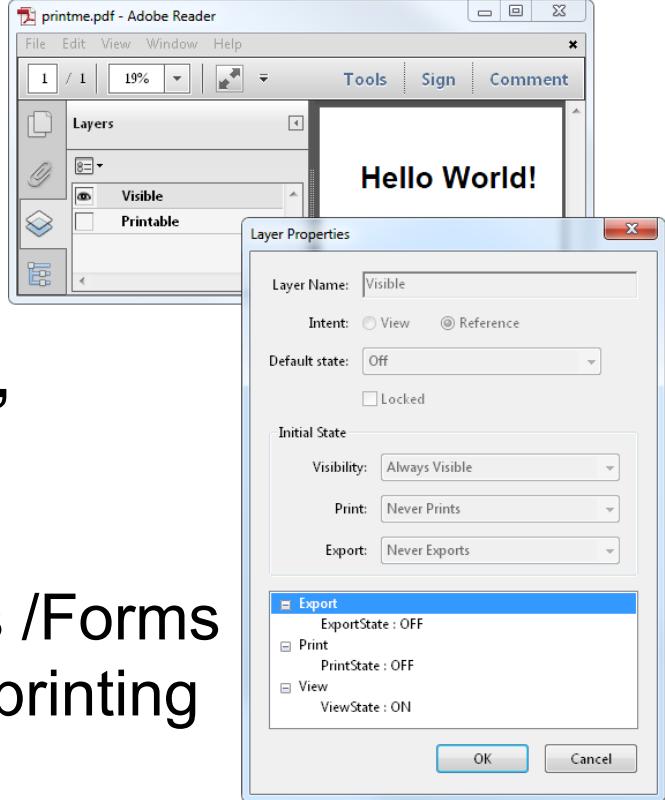


What you see is not what you'll get...

PDF Layers 1/2

“Optional Content Configuration”

- principles
 - define layered content via various /Forms
 - enable/disable layers on viewing/printing
- no warning when printing
- “you can see the preview!”
 - bypass preview by keeping page 1 unchanged
 - just do a minor change in the file

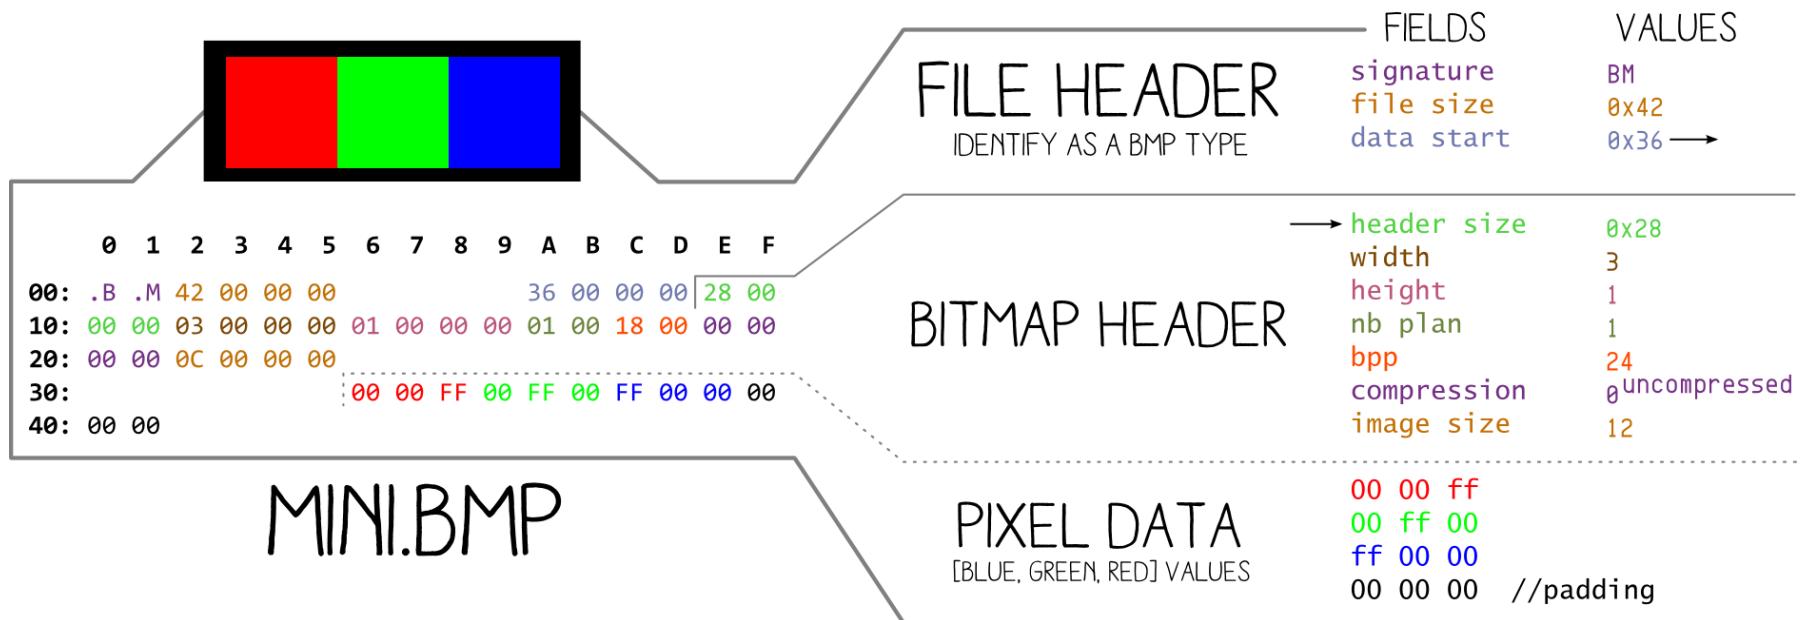


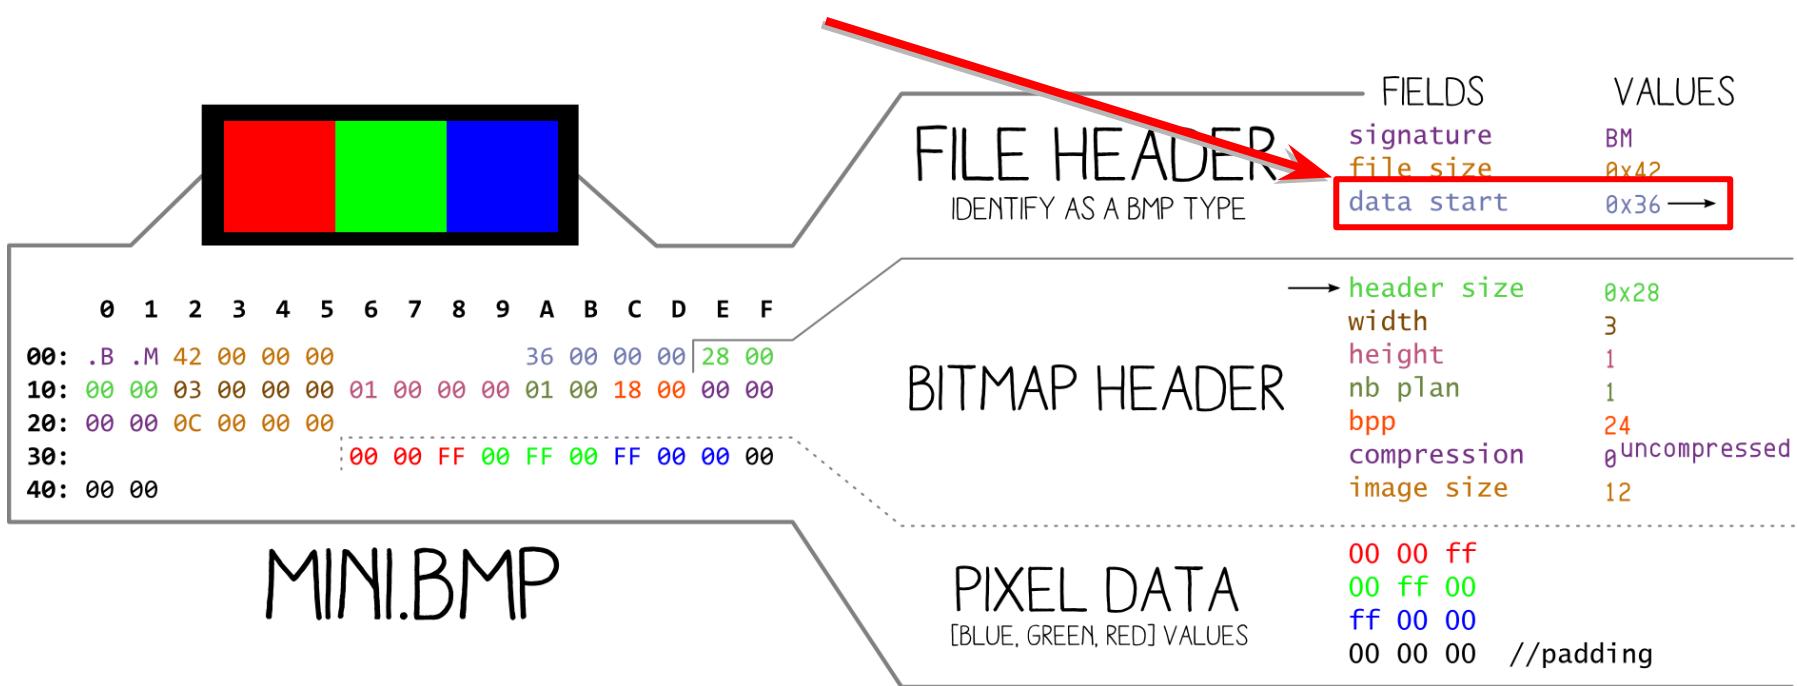
PDF Layers 2/2

- it's Adobe only
 - what's displayed varies with readers
 - could be hidden via previous schizophrenic trick
- it was in the specs all along
 - very rarely used
 - can be abused with no warning

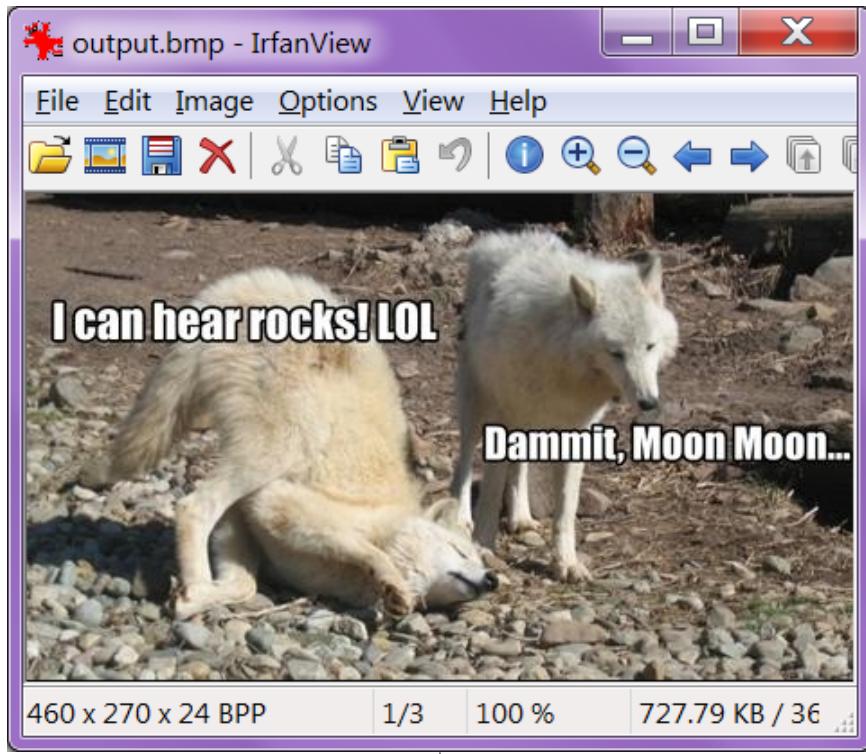
BITMAP / DEVICE INDEPENDENT BITMAP

SUBTYPE TYPE





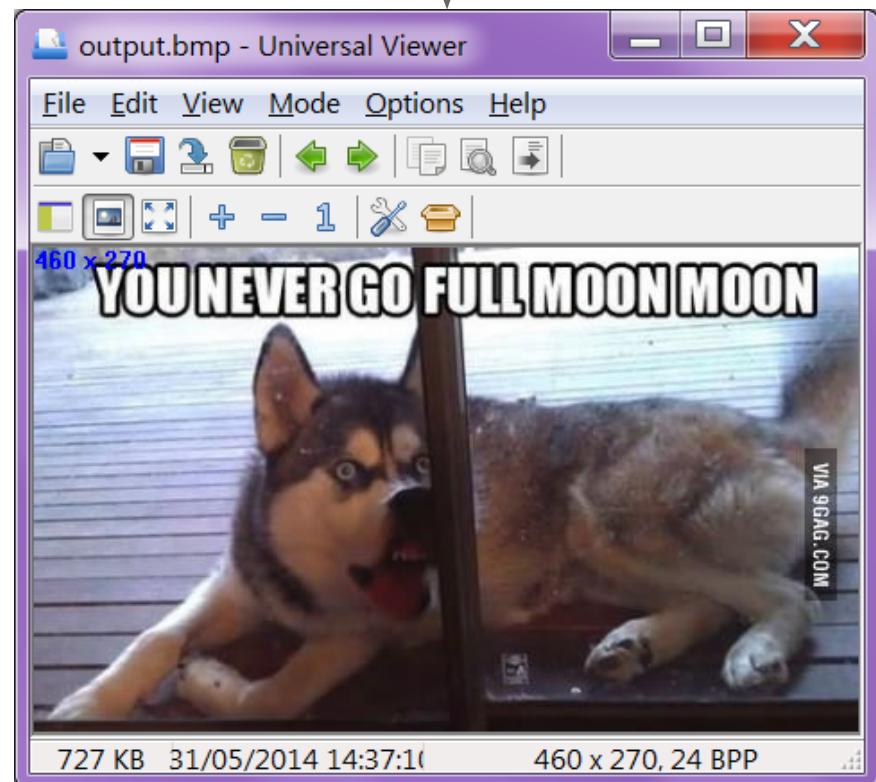
A pointer to some information that usually comes next...
 What could go wrong...



getting data
via the pointer
(standard)

BMP Trick #1: ignoring the data pointer

getting data right after
the header



Trick #2:

Run-Length Encoding

BMP RLE trick

RLE structure (each box is 1 byte)

Length >0	Palette Index (color)			
Length 0	RAW Length >2	Palette Index (color)	...	Palette Index (color)
Length 0	End of Line 0			
Length 0	End of Bitmap 1			
Length 0	Move Cursor 2	X offset		Y offset

BMP RLE trick

If you just skip pixels, what is their color?

Length 0	End of Line 0
--------------------	-------------------------

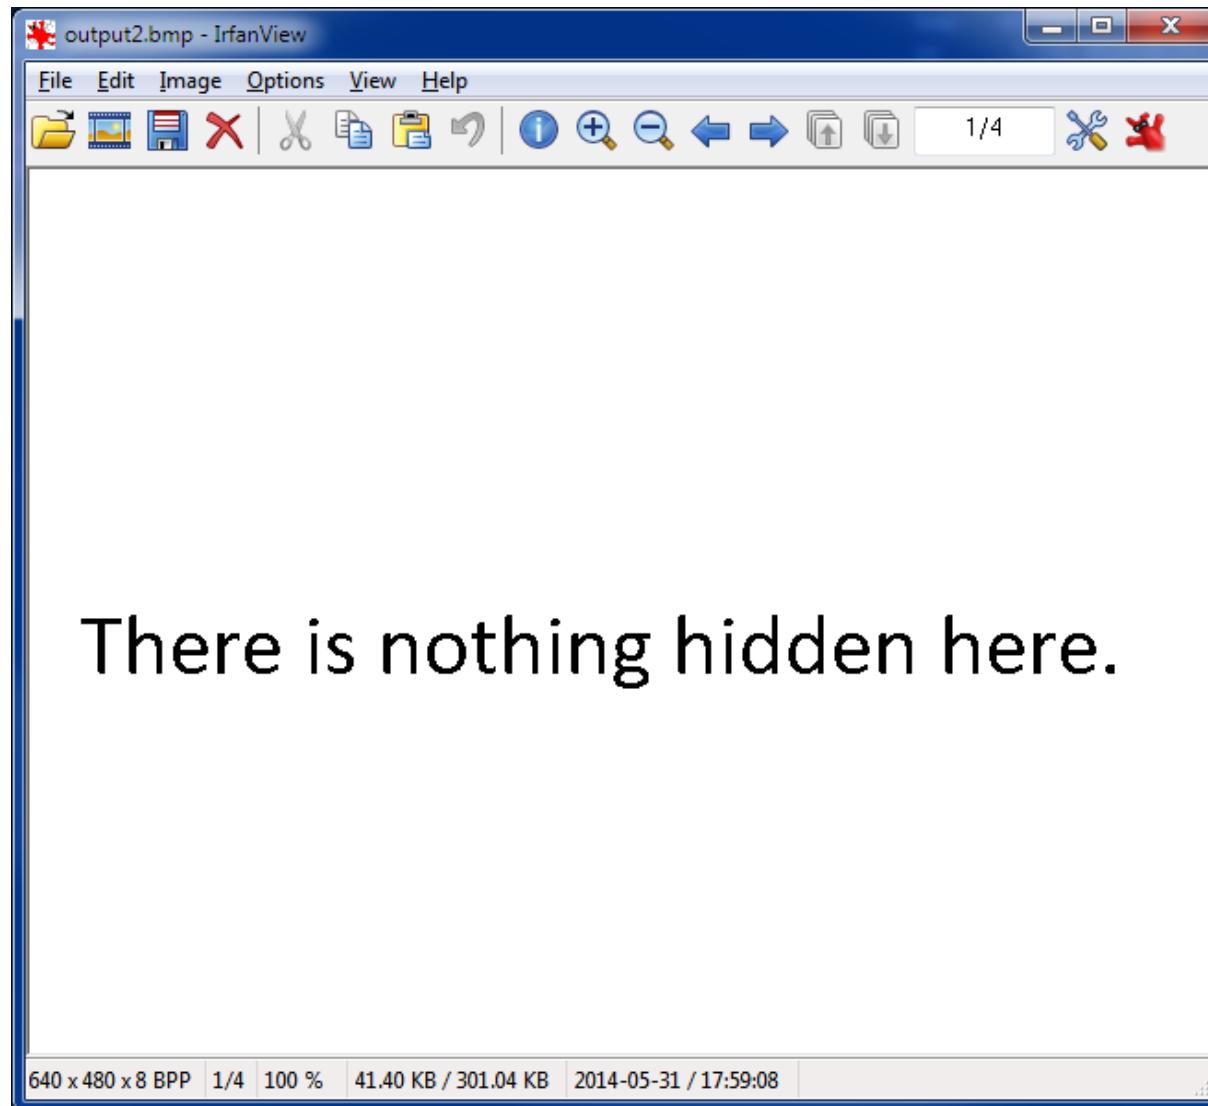
Length 0	End of Bitmap 1
--------------------	---------------------------

Length 0	Move Cursor 2	X offset	Y offset
--------------------	-------------------------	-----------------	-----------------

Option 1

The missing data will be filled with **background color**.

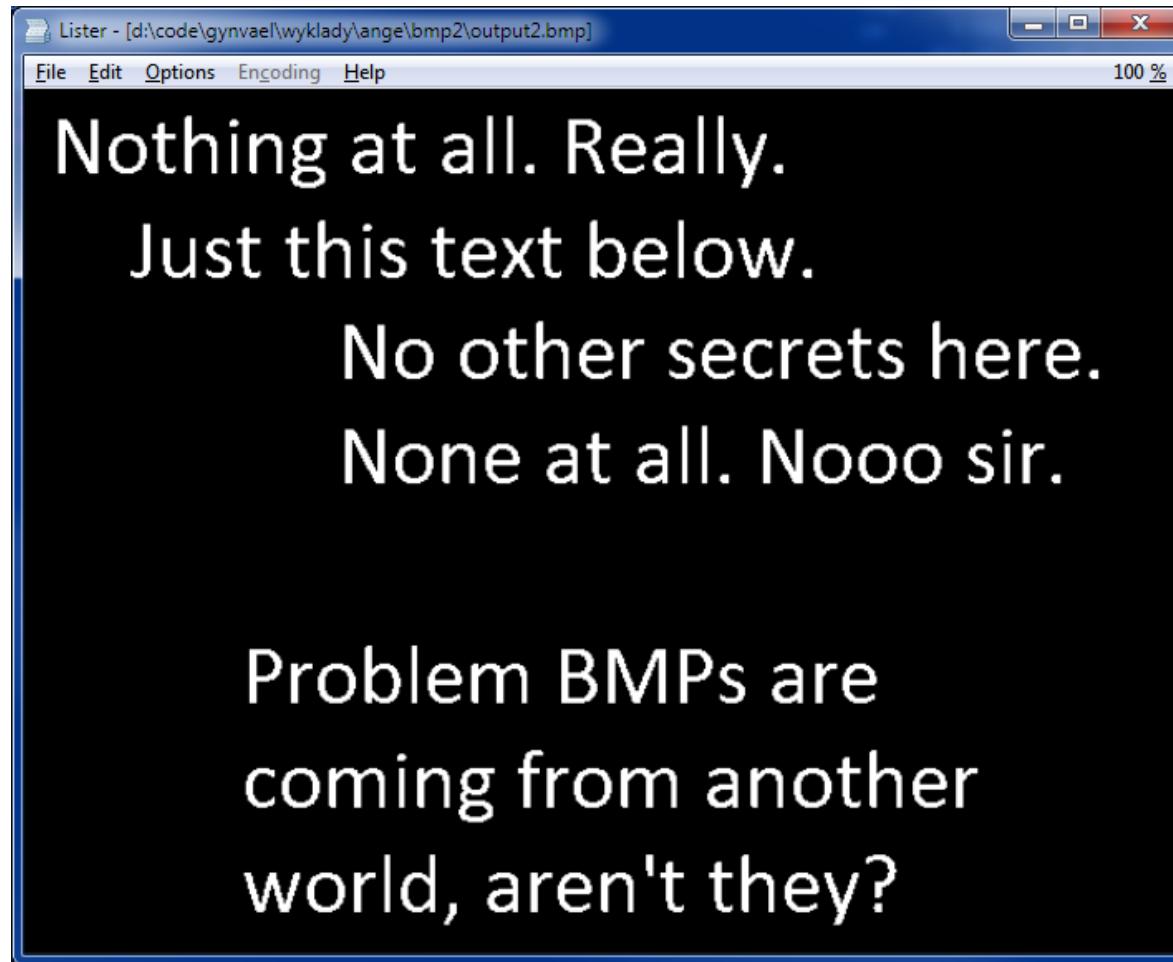
(palette index 0)



There is nothing hidden here.

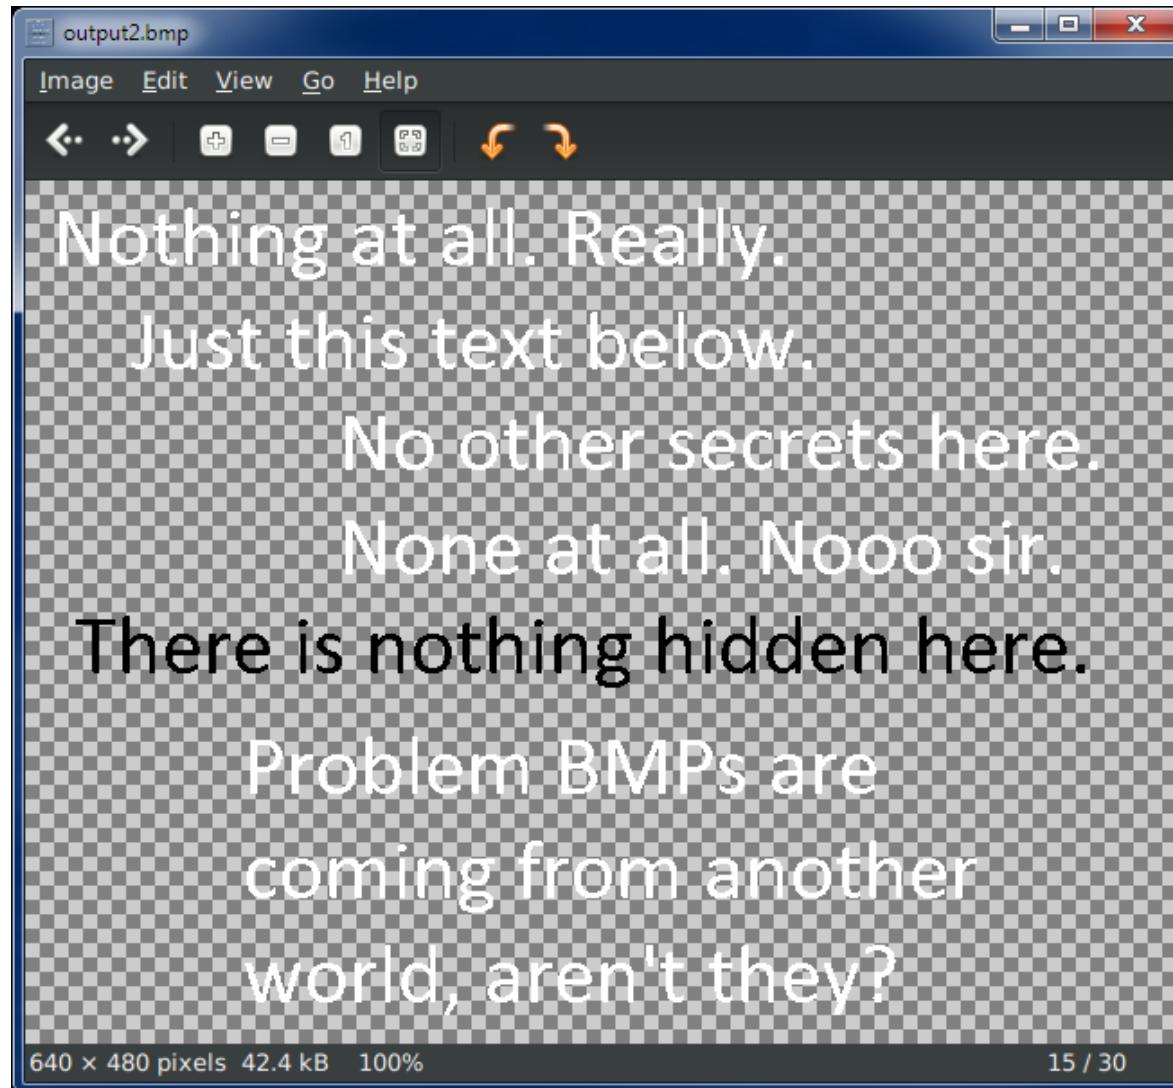
Option 2

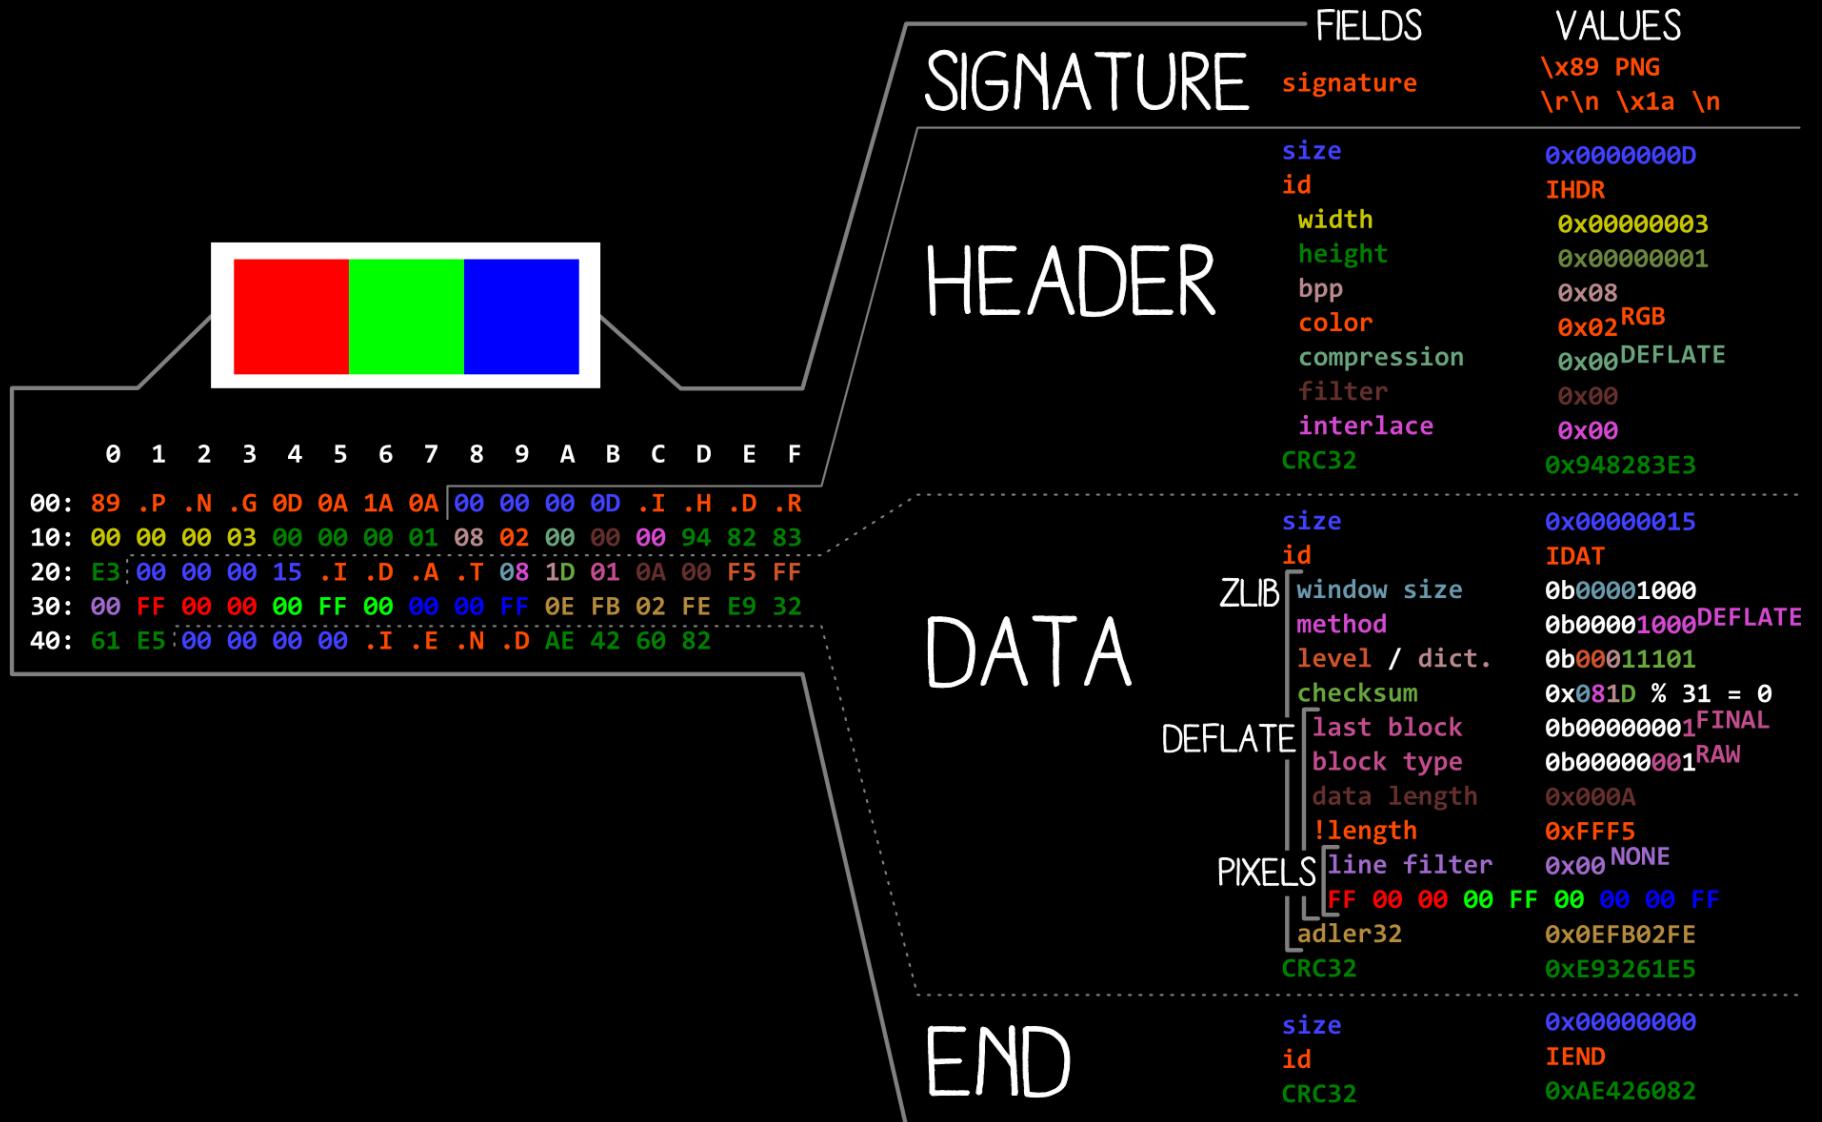
The missing data will be **black**.



Option 3

The missing data will be **transparent**.





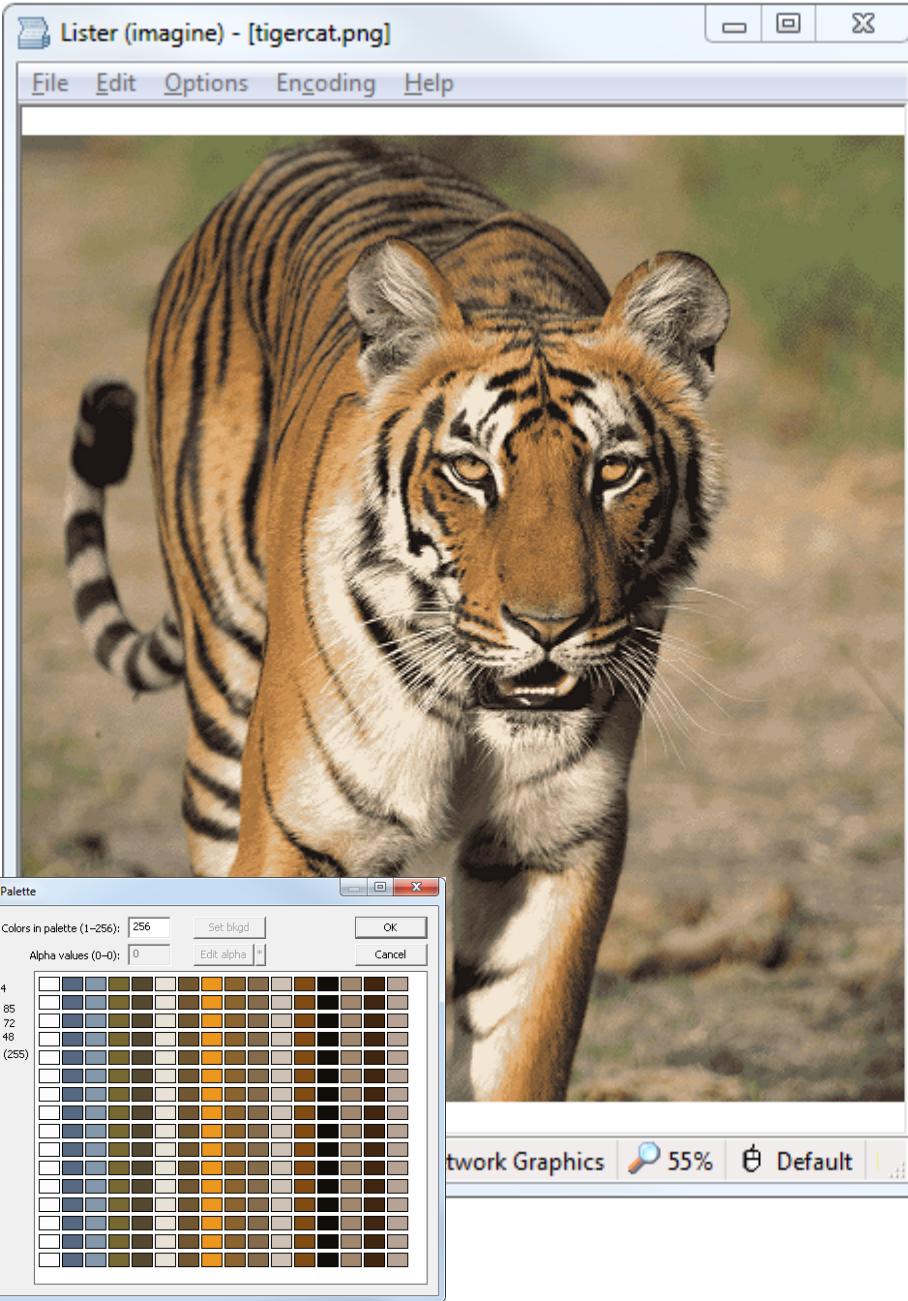
Portable Network Graphics

Combined data + 2 palettes

Same data chunk combining 2 images via 2 palettes

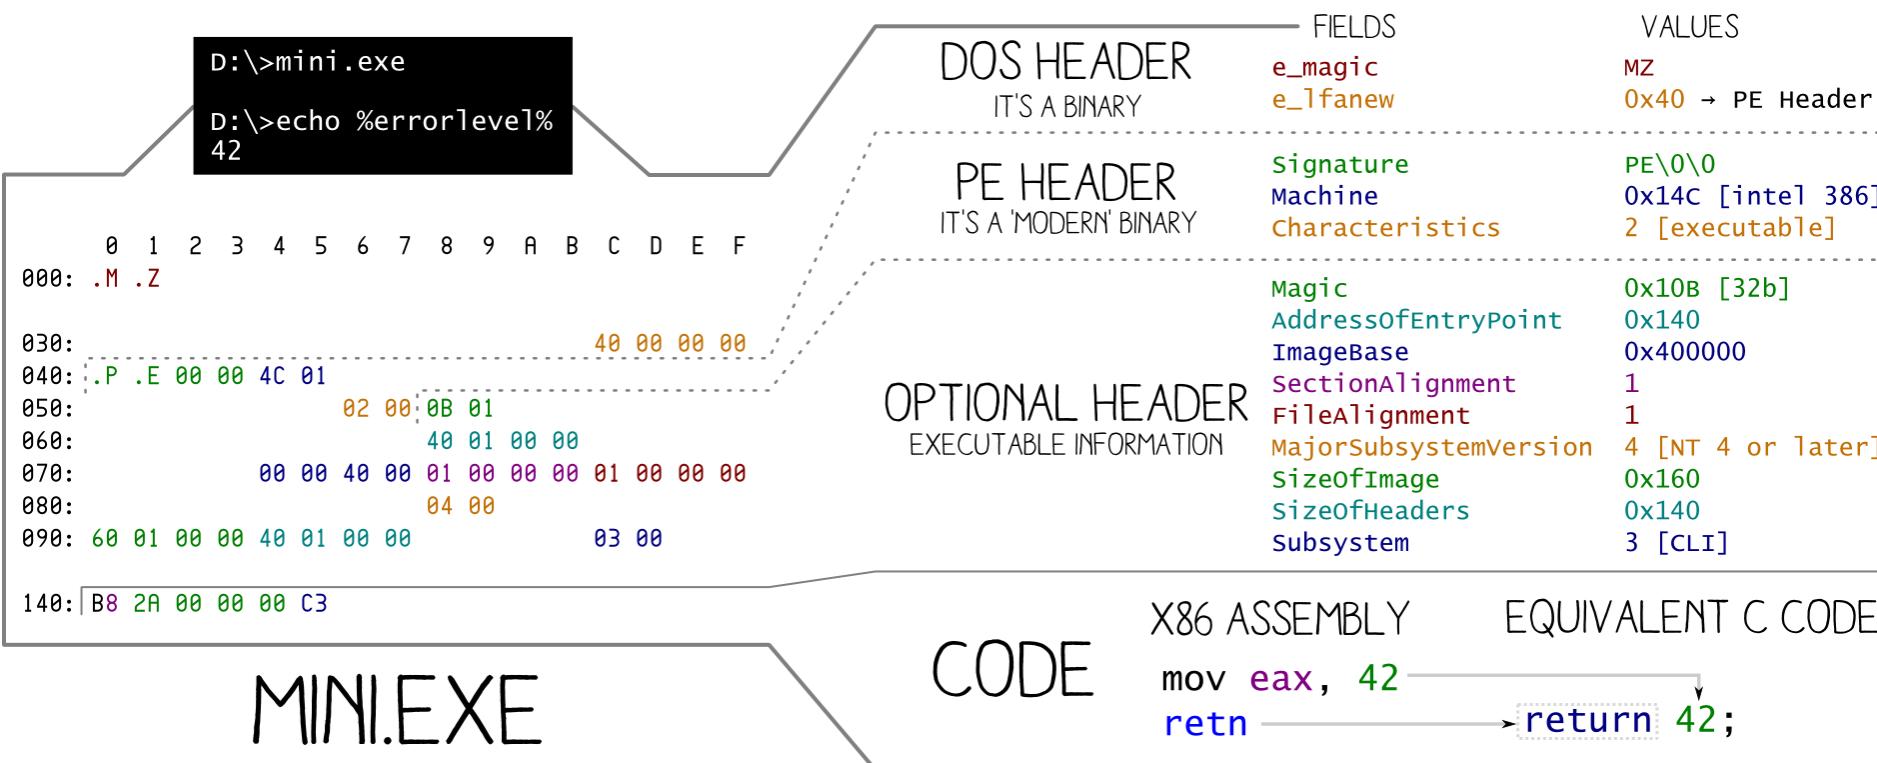
cute PoC by [@reversity](#)

“There shall not be more than one PLTE chunk”



Different images depending on which PLTE chunk is used

PORTABLE EXECUTABLE



the PE Loader

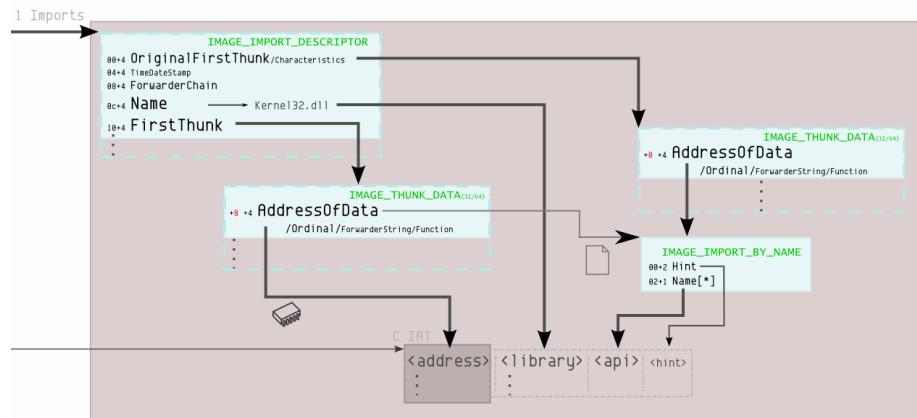
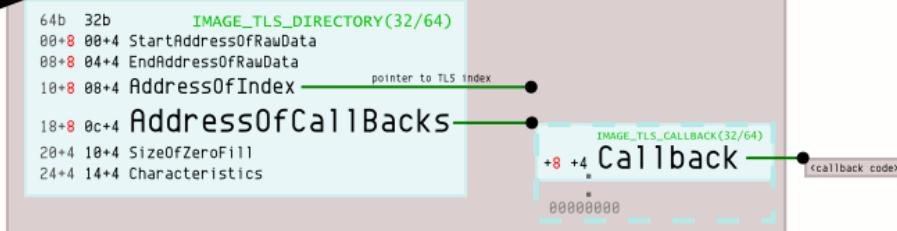
PE = complex + badly documented

- fail or fool external tools ? too easy...
- fooling Windows is much harder:
 - Windows' loader usually closes holes
⇒ older PEs just not working anymore

PE Trick #1

Data directory loading order

9 Thread Local Storage



TLS
AddressOfIndex →

Imports
descriptor #1
descriptor #2
descriptor #3

Pointing TLS' AddressOfIndex to an Import descriptor

```
C:\ demo AoI OS Detection  
C>ver  
Microsoft Windows [V  
C>tlS_aoI0SDET.exe  
* TLS AoI => W7
```

TLS
AddressOfIndex →

Imports
descriptor #1

descriptor #3

loaded first

W7: TLS is loaded first ⇒ AoI's address set to 0
⇒ Imports descriptors's sequence is truncated before loading

```
c:\> TLS AoI on imports  
D>ver  
Microsoft Windows XP [U]  
D>tls_aoiOSDET.exe  
* TLS AoI => XP
```

TLS
AddressOfIndex →

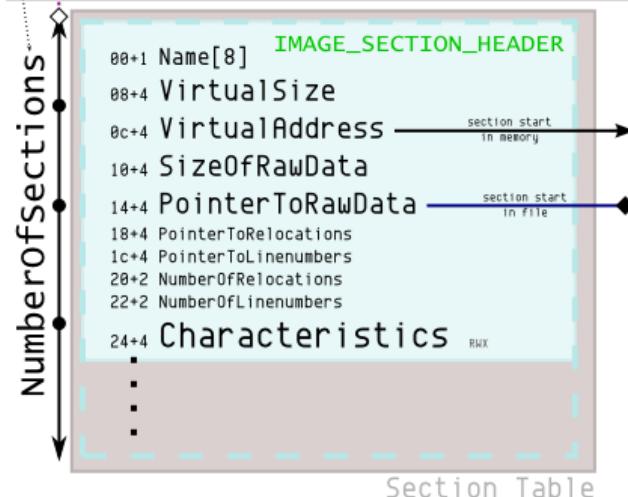
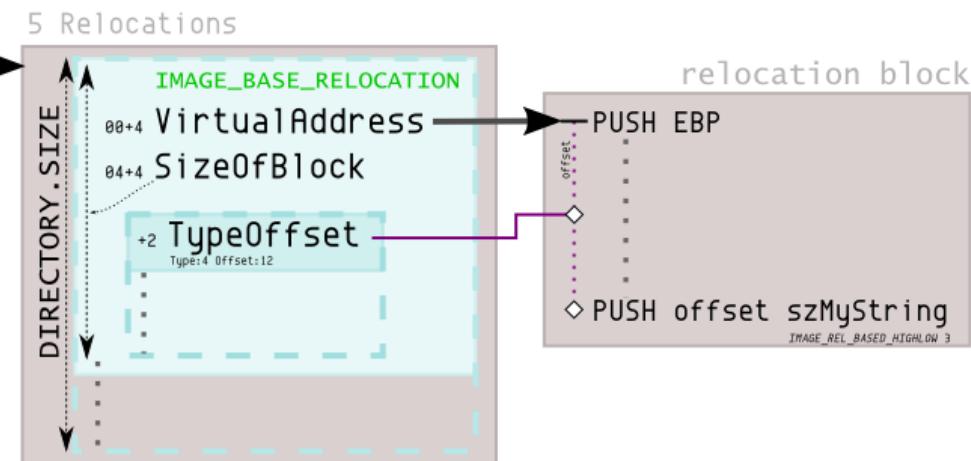
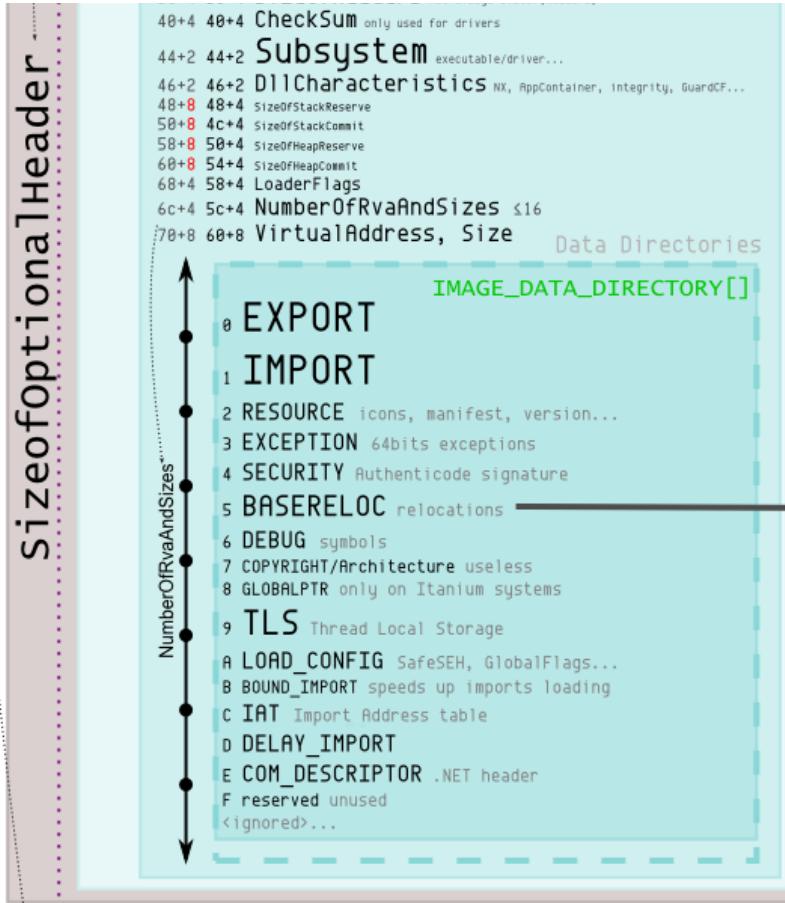
Imports
descriptor #1
descriptor #2
descriptor #3

loaded first

XP: Imports are loaded first - all descriptors are parsed
TLS is then parsed - descriptors are not relevant anymore

PE Trick #2

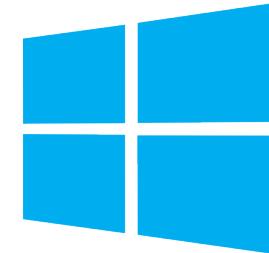
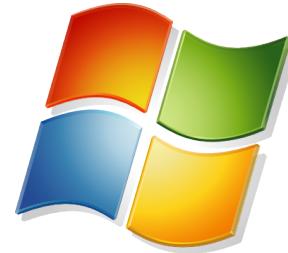
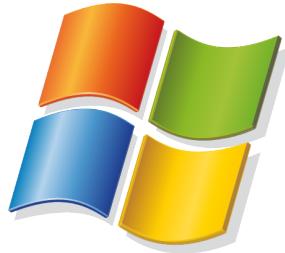
Relocations



Relocations:

patching absolute addresses
to solve address space conflicts

Relocations types



Type 4
HIGH_ADJ

--

--



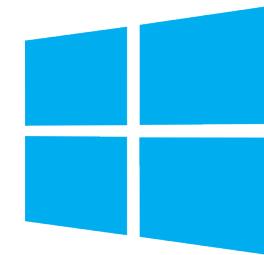
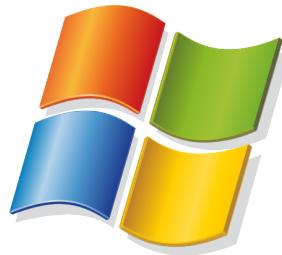
Type 9
MIPS_JMPADDR16
IA64_IMM64
MACHINE_SPEC_9

32 bit

64 bit



Relocations on relocations



Type 4
HIGH_ADJ

--

--

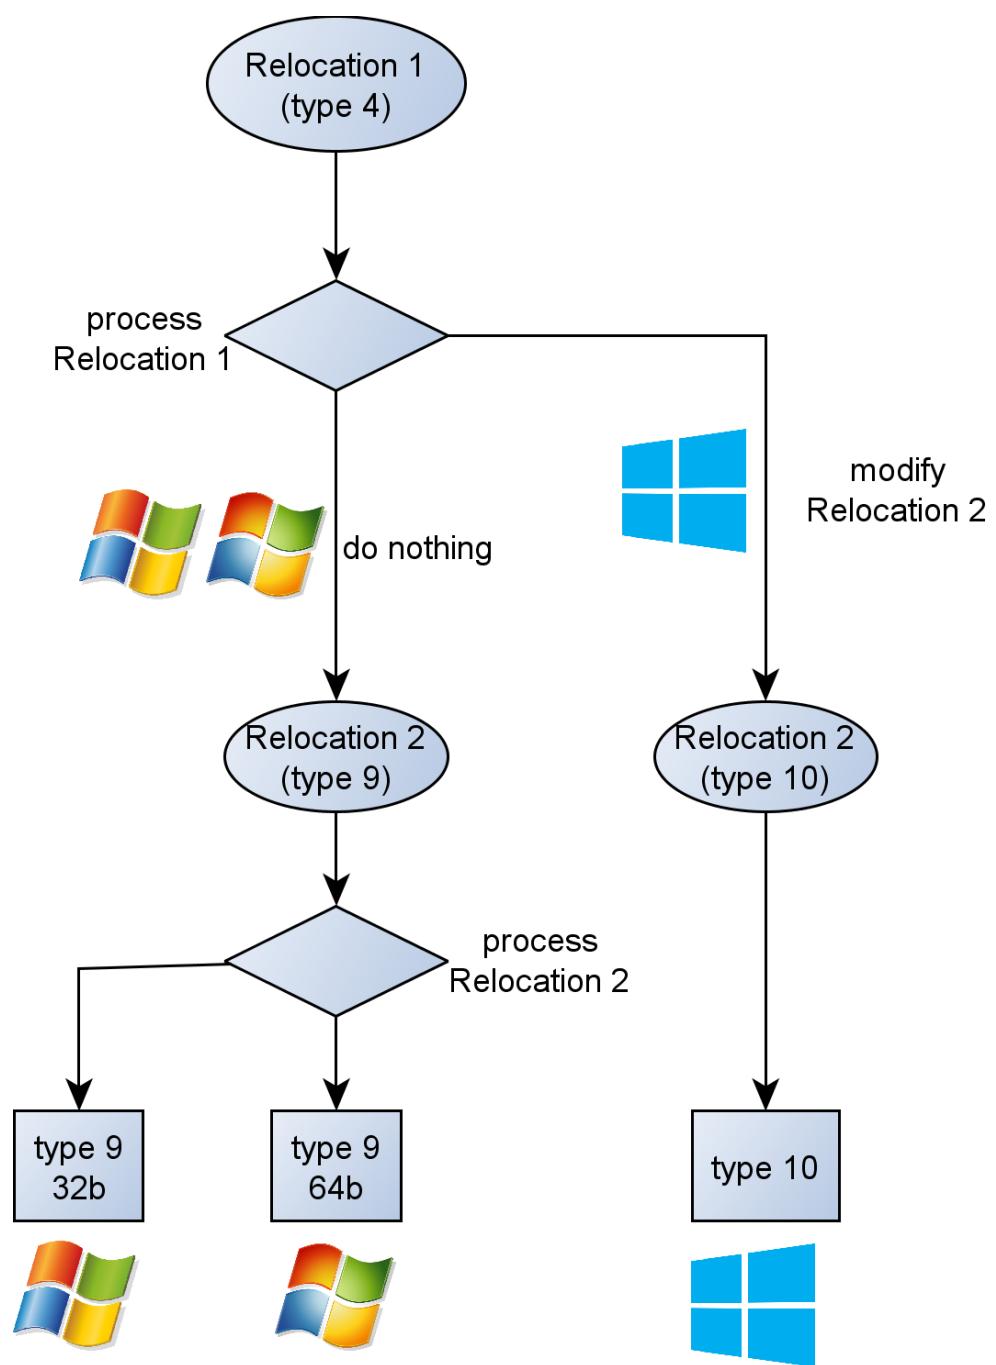
Type 9
MIPS_JMPADDR16
IA64_IMM64
MACHINE_SPEC_9

32 bit

64 bit

Type 10
DIR64





Relocation-based PE Schizophren

FOLLOW US INTO THE RABBIT HOLE

```
./hello  
dlrow ,olleH
```



IDA uses sections

Kernel uses segments

```
...  
lic main  
t near DATA WORD _startWithRe  
b rbp, rbp  
    edi, offset hello_2 : "Hello, World!"  
l _putchar  
r rbp  
  
'Hello, World\n'
```

CONFIDENCE 2013

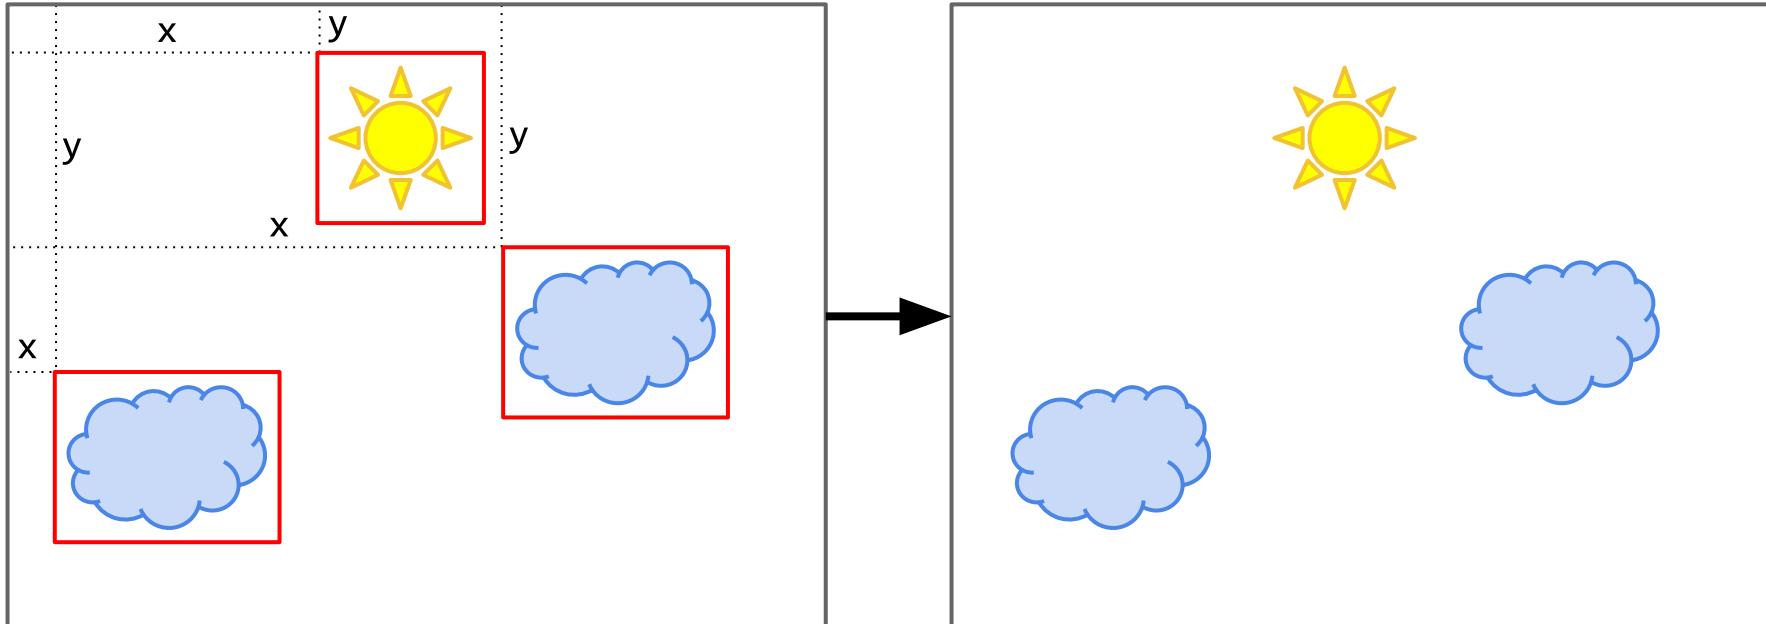


Julian Bangert, Sergey Bratus -- ELF Eccentricities
<https://www.youtube.com/watch?v=4LU6N6THh2U>

GIF

GIF

A GIF is made of blocks.
if no animation speed is defined,
they should all be displayed at once.

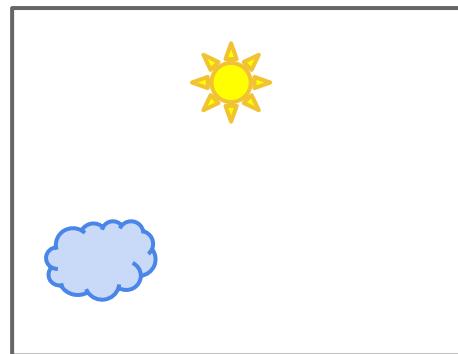


GIF

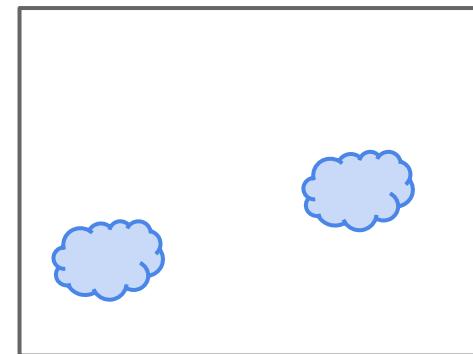
If a frame speed is defined, then:
first block = background
next blocks = animation frames



Background
(from block 1)



Frame 1
(with block 2)



Frame 2
(with block 3)

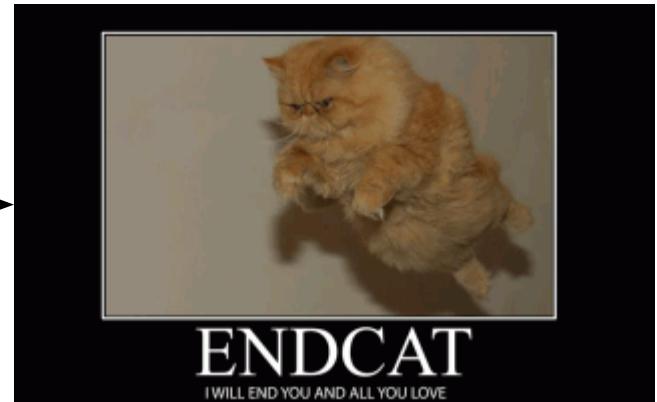
GIF



Frame 1



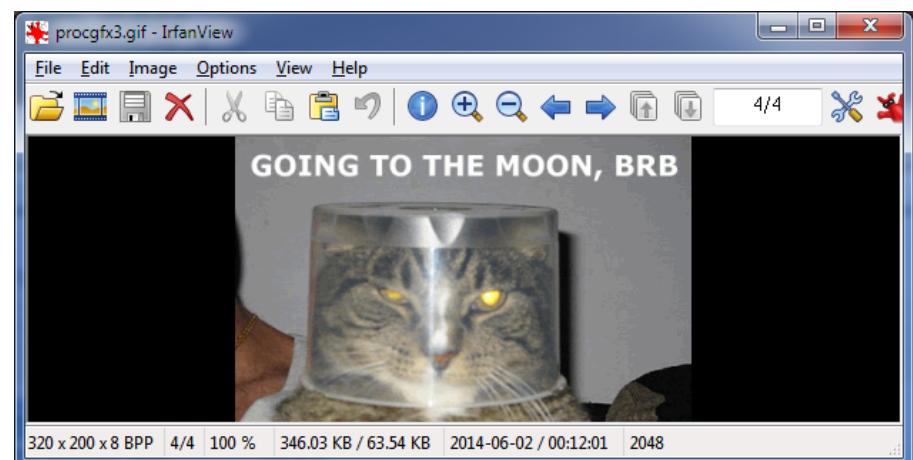
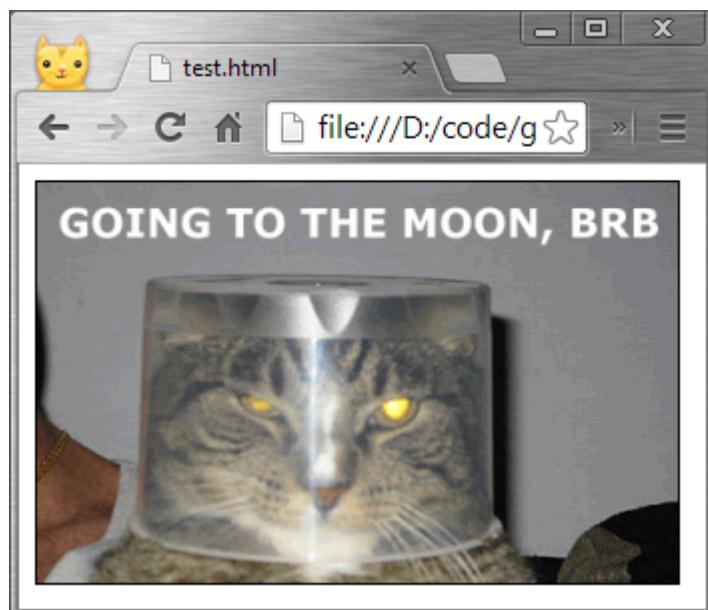
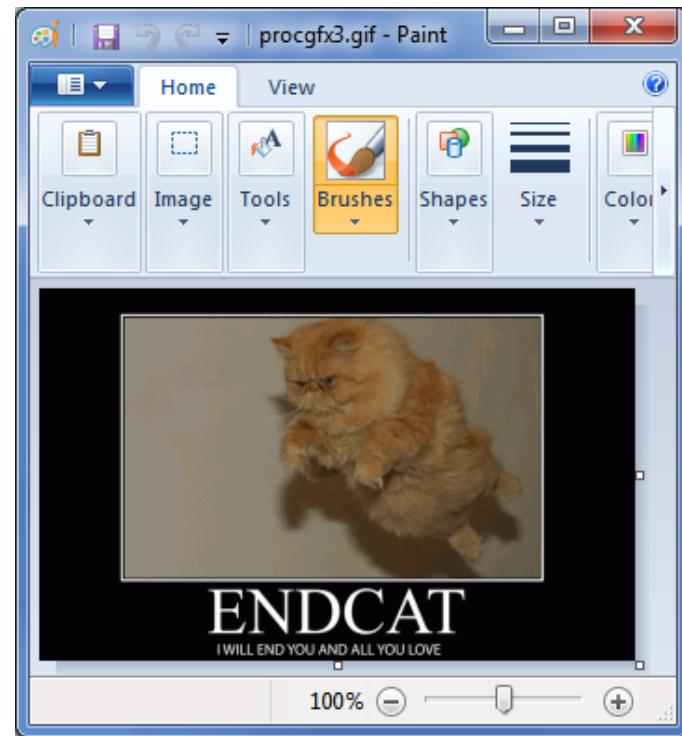
Frames 2-10001
1x1 px



Frame 10002

1 complete pic + 10.000 pixels + 1 complete pic

Displaying all blocks at once (standard)



Forcing animation (even if no frame speed is defined)

same-tool schizophrenia

1 file + 1 tool = 2 behaviors
(in different sub-components)

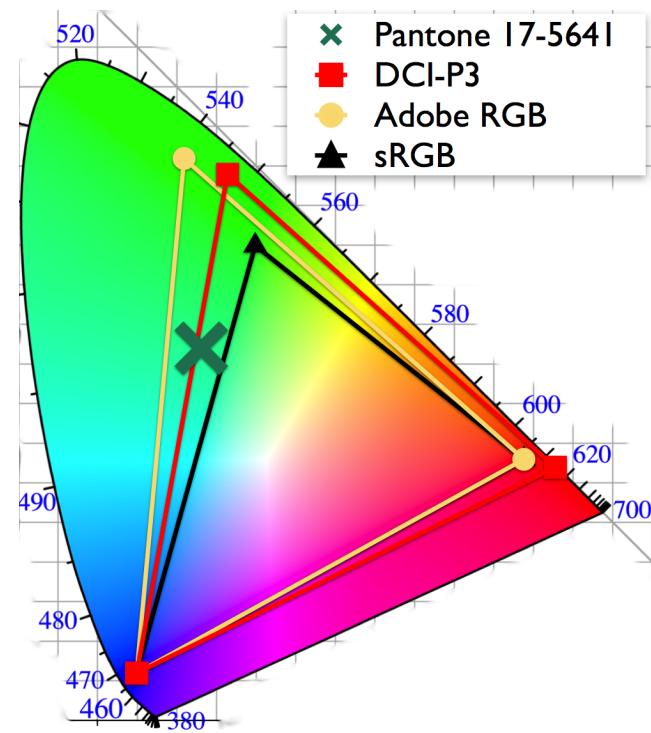
Because it was too simple...

- WinRar: viewing \Leftrightarrow extracting
 - opening/failing
 - opening/'nothing'
- Adobe: viewing \Leftrightarrow printing
 - well, it's a *feature*

Failures & Ideas

Failures & Ideas

- screen ⇔ printer
 - embedded color profiles?
- JPG
 - IrfanView vs the world
- Video
 - FLV: early data pointer, like BMP
 - PoC: video fails but plays sound



PNG

Various ancillary chunks (rendering level)

- partially supported:
 - gamma
 - transparency (for palettes)
- never supported?
 - significant bits
 - chromacities
- always supported?
 - physical size

Conclusion

A photograph of a park scene. A paved path leads from the foreground towards a grassy area. On the right side of the path, there is a grassy slope where a person is walking. The sky is overcast. The text "Design" is overlaid on the paved path, and "User Experience" is overlaid on the grassy slope.

Design

User Experience

We tend to take our own shortcuts.

Conclusion

- such a mess
 - specs are messy
 - unclear
 - historical reasons
 - parsers don't even respect them
(particularly when there is an easy shortcut)
 - official tools “forced” to be tolerant
 - They're even trying to repair corrupted files (!)
- no CVE/blaming for parsing errors?
 - no security bug if no crash or exploit :(

Schizophrenia symptoms

- different parsing (seeing different data)
 - BMP: ignoring data pointer
 - ZIP: different parsing algorithm & directions
 - PE: different data directory loading order
 - PDF: different trailer parsing
- different interpretation (same data)
 - GIF: ignoring animation speed
 - BMP RLE: using different default color
 - PE: different relocations implementation
 - PNG: using different palette
 - PDF: conditional layers

ACK

@gynvael

@reversity @travisgoodspeed @sergeybratus
qkumba @internot @pdfkungfoo

@j00ru ise ds vx, Mulander
Felix Groebert, Salvation

@angealbertini
corkami.com

SCHIZO
PHRENIC

FOLIES

